

Systems Design

Carlos Cuevas
Brooklyn College

Typical Software Engineer Interview

- Data Structures & Algorithms
- Practical Coding Interview
- Behavioral
- Systems Design

Systems Design Interview

- Design a system that accomplishes x
 - Chat System
 - URL Shortener
 - Youtube
 - Many Others

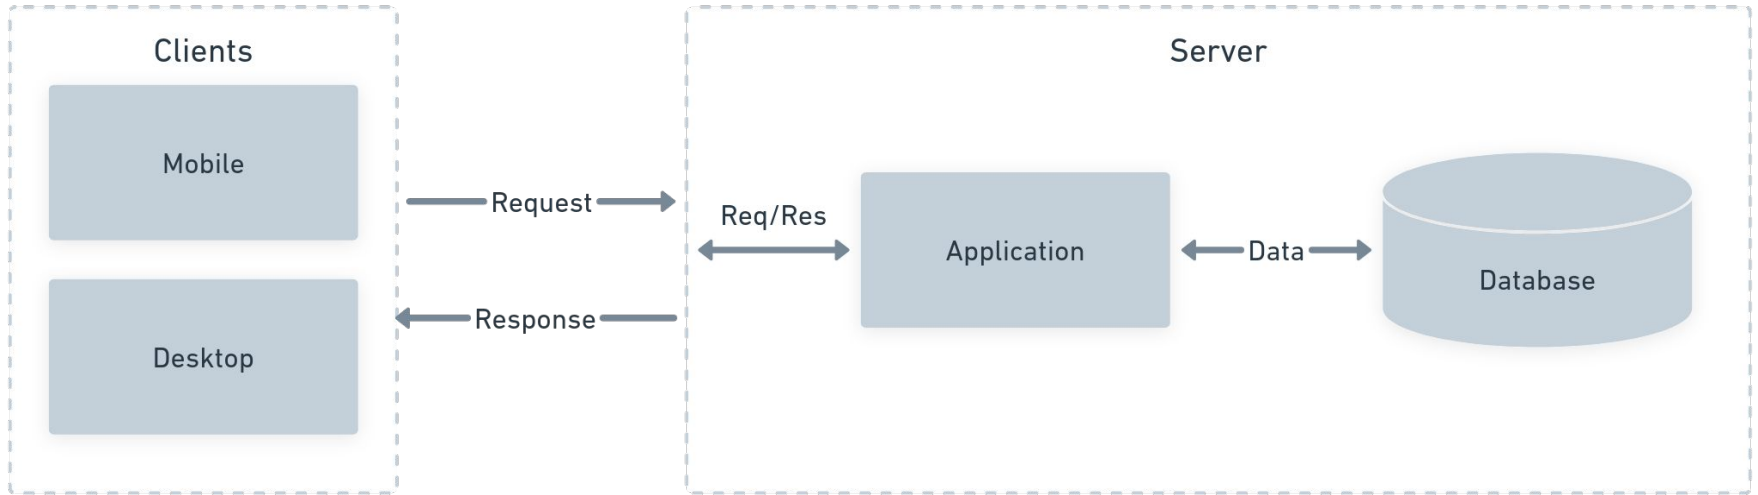
Systems Design Interview

- It's about the journey, not the destination.
 - There isn't one right answer.
 - Communication is as (if not more) important than your design.
- Above all, do not jump straight into designing. Ask **a lot** of questions first.

Building Blocks

- Many systems have overlap
- Learn the common components of many large scale systems

Client - Server Model



What exactly is a Server?

- Hardware
 - A physical* computer
 - on which **server software** is running
- Software
 - A process that is waiting for incoming requests and responding to those requests (i.e. serving)

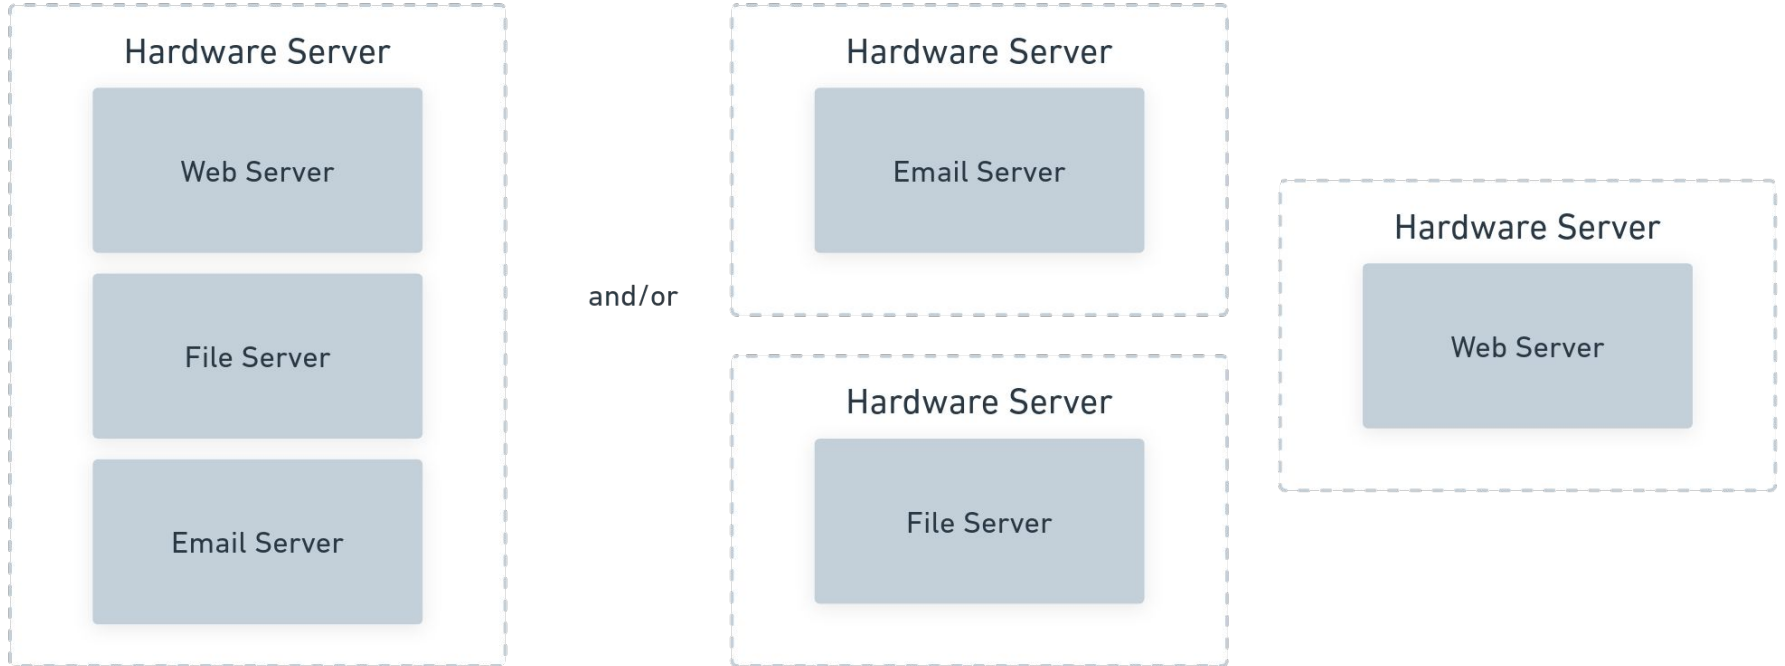
* Could be:

- ❖ Physical Computer
- ❖ Virtual Machine / Container
- ❖ Cloud Service

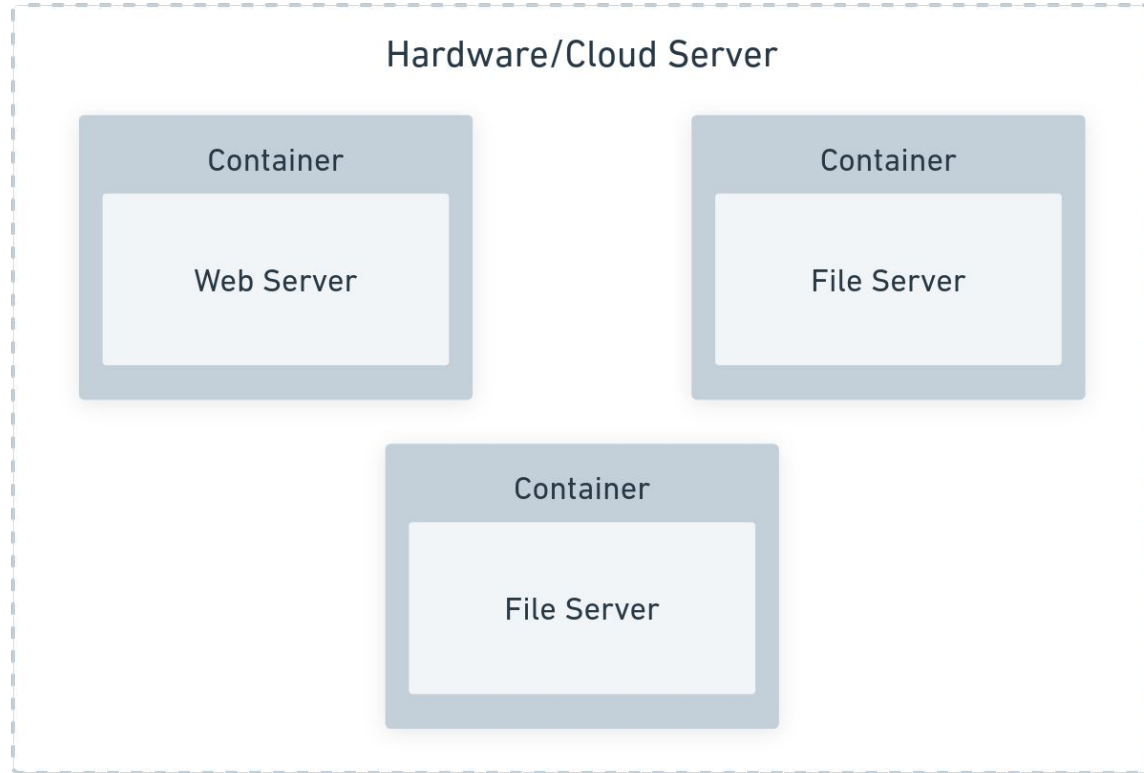
Types of (Software) Servers

- **Web**
 - Apache
- **File**
 - FTP
- **Email**
 - SMTP
- **Database**
 - MySQL
- **Many many more**

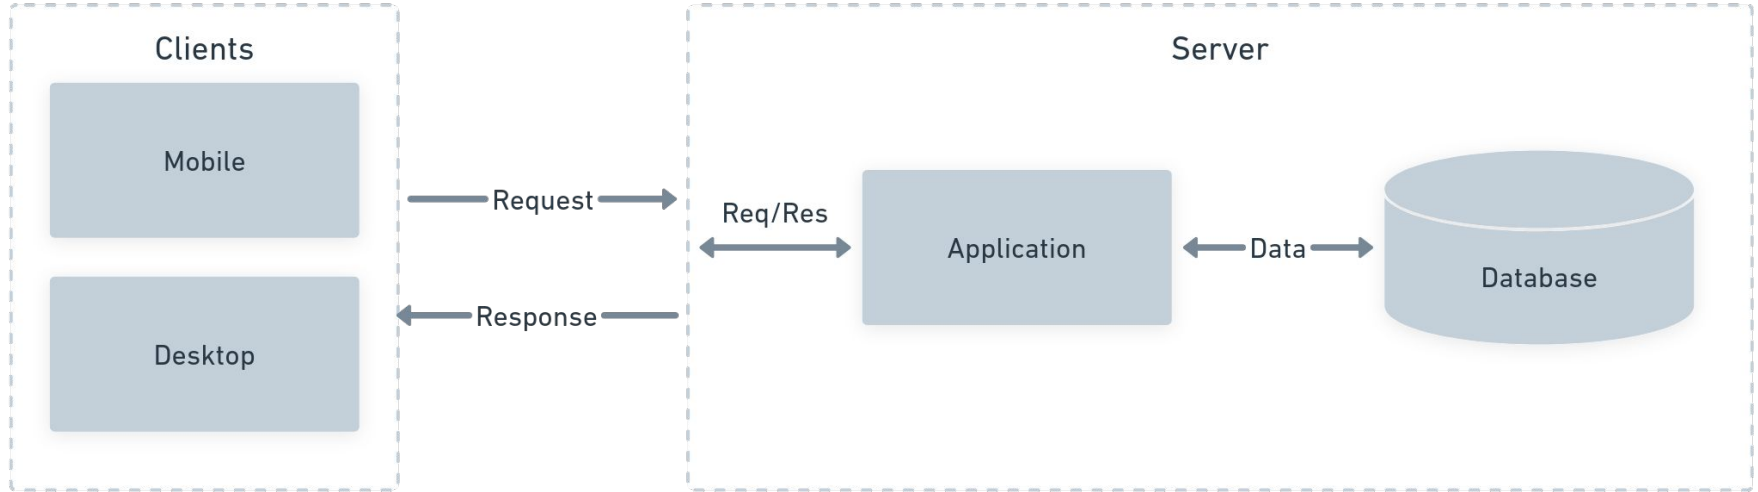
Hardware & Software



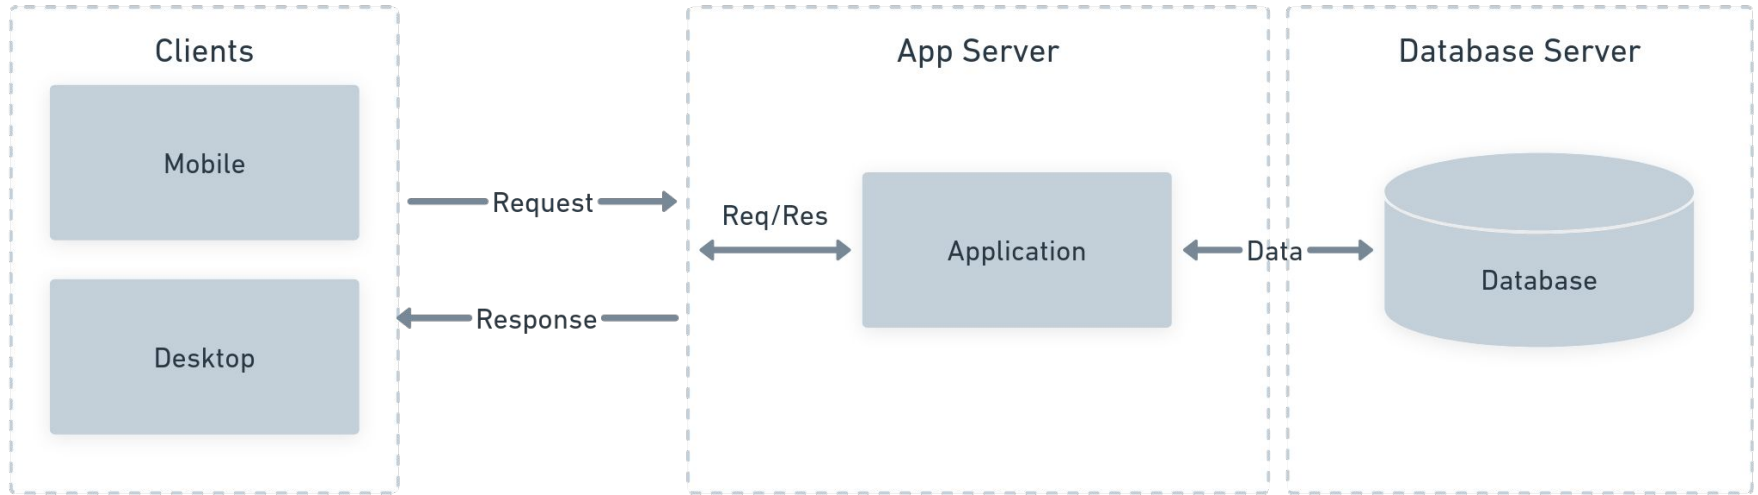
Containerization




Single Server



Multi-Server Setup



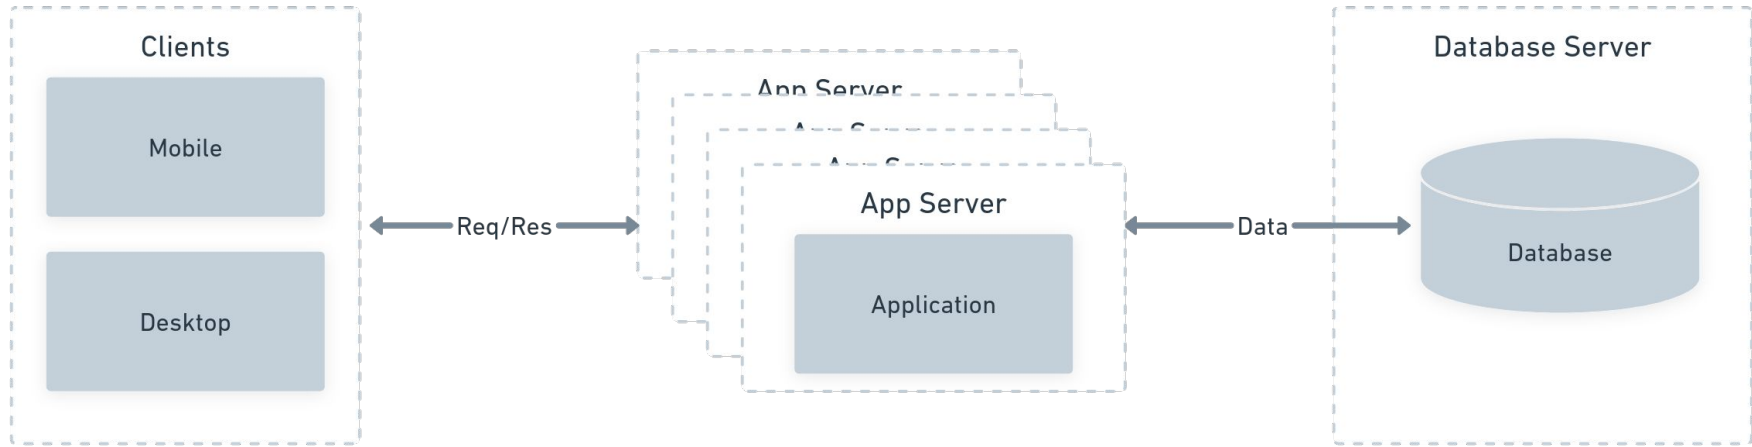
Horizontal vs Vertical Scaling



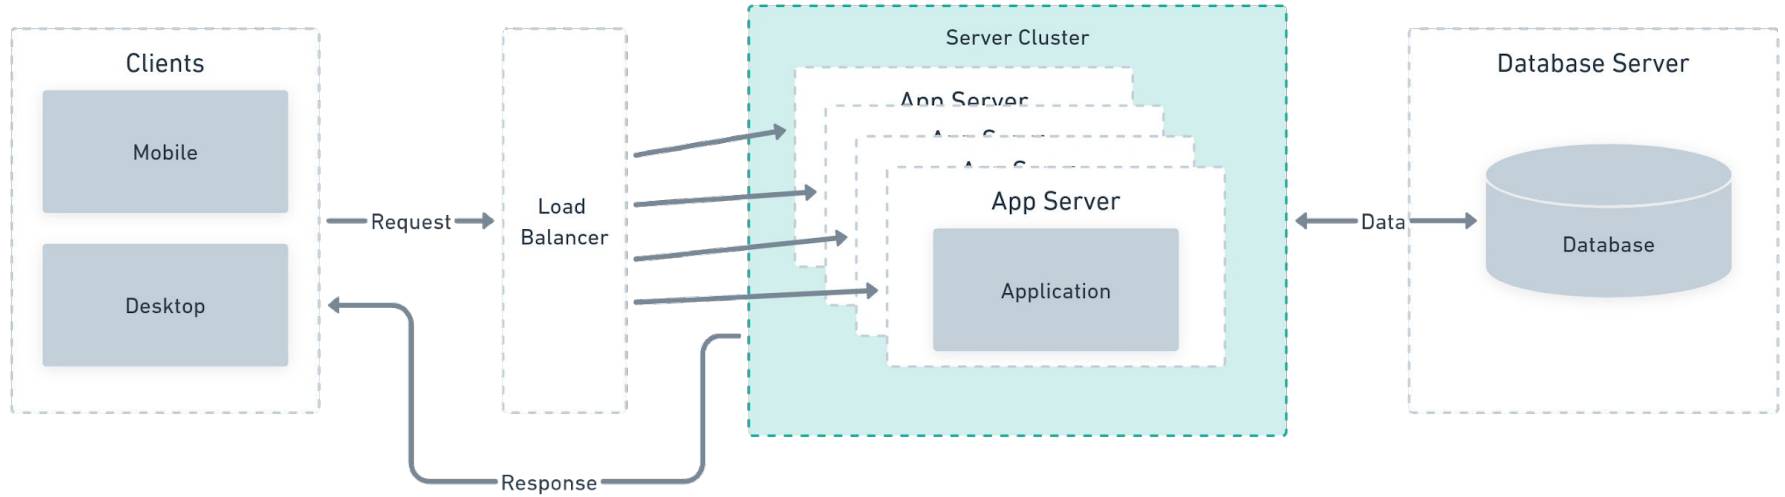
Horizontal vs Vertical Scaling

- There's a limit on how much you can vertically scale
 - After you've maxed out the most powerful CPU in existence, then what?
- Horizontal Scaling is, essentially, infinite
 - Just add more instances and distribute the load

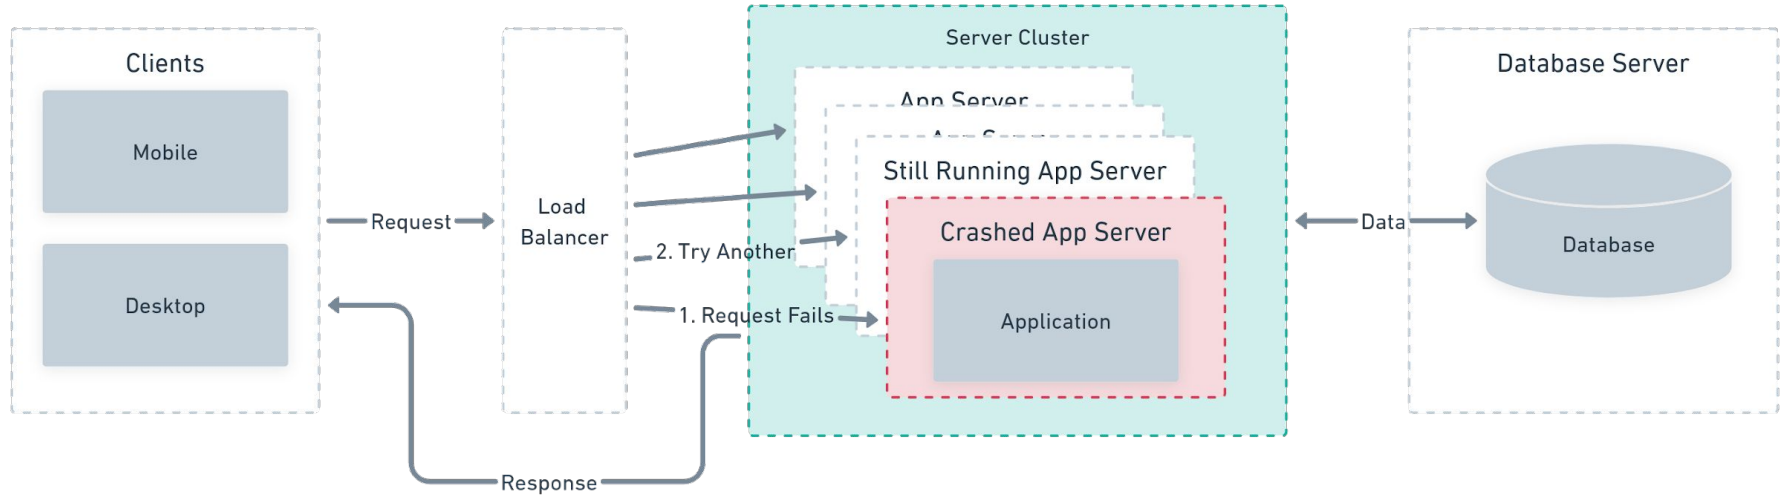
Horizontal Scaling



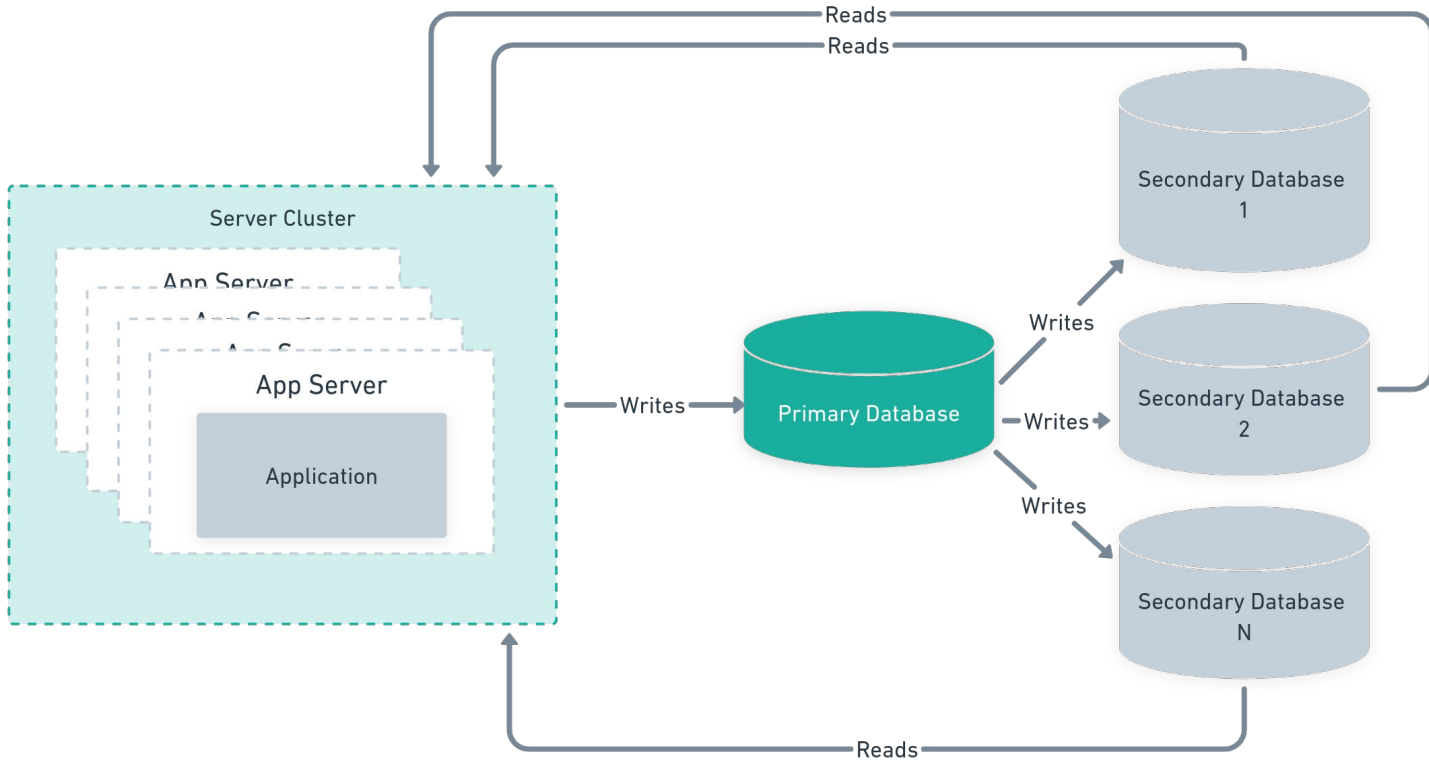
Load Balancing



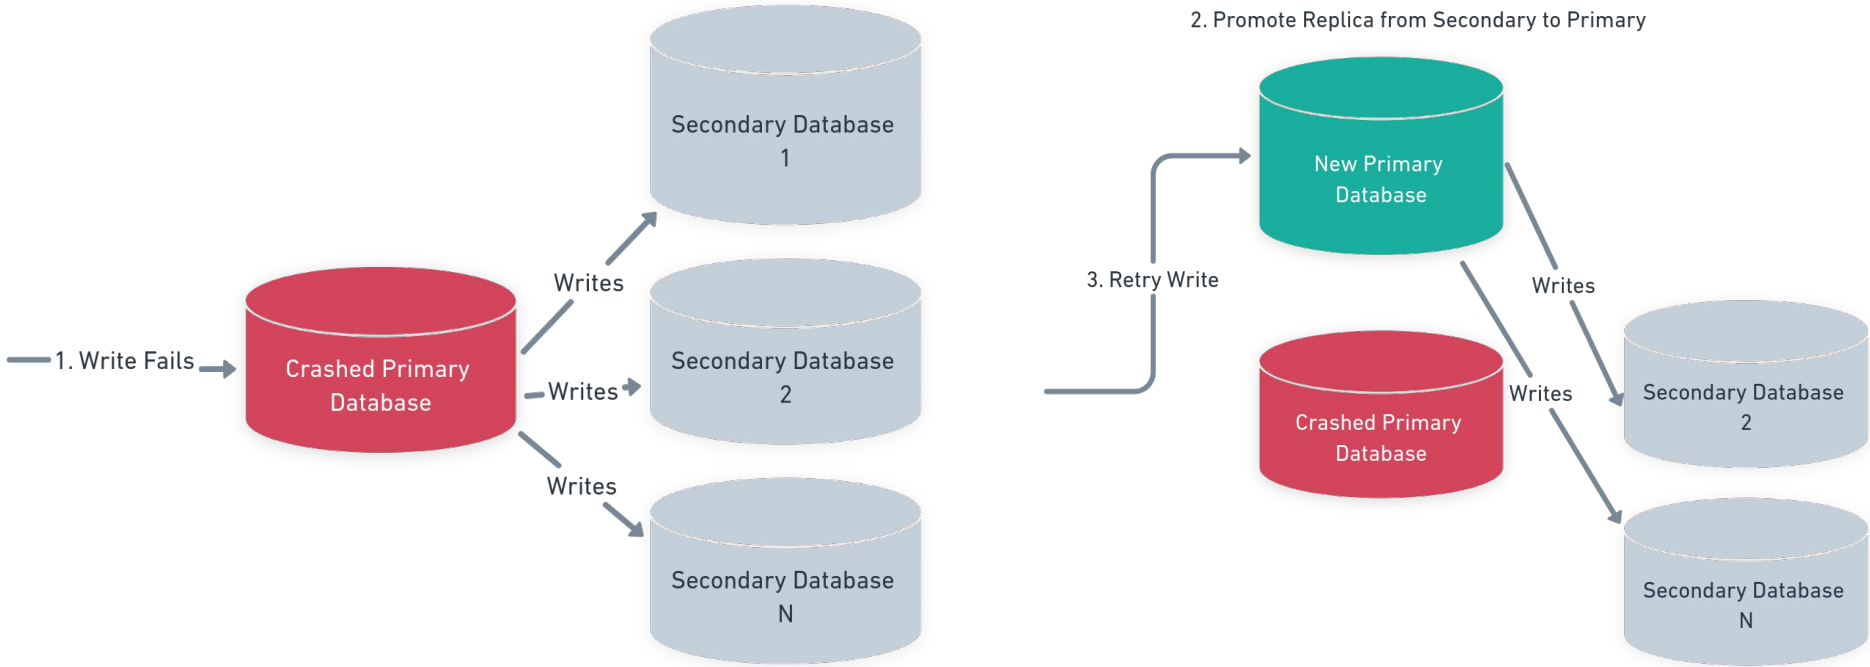
Horizontal Scaling + Load Balancing = Availability



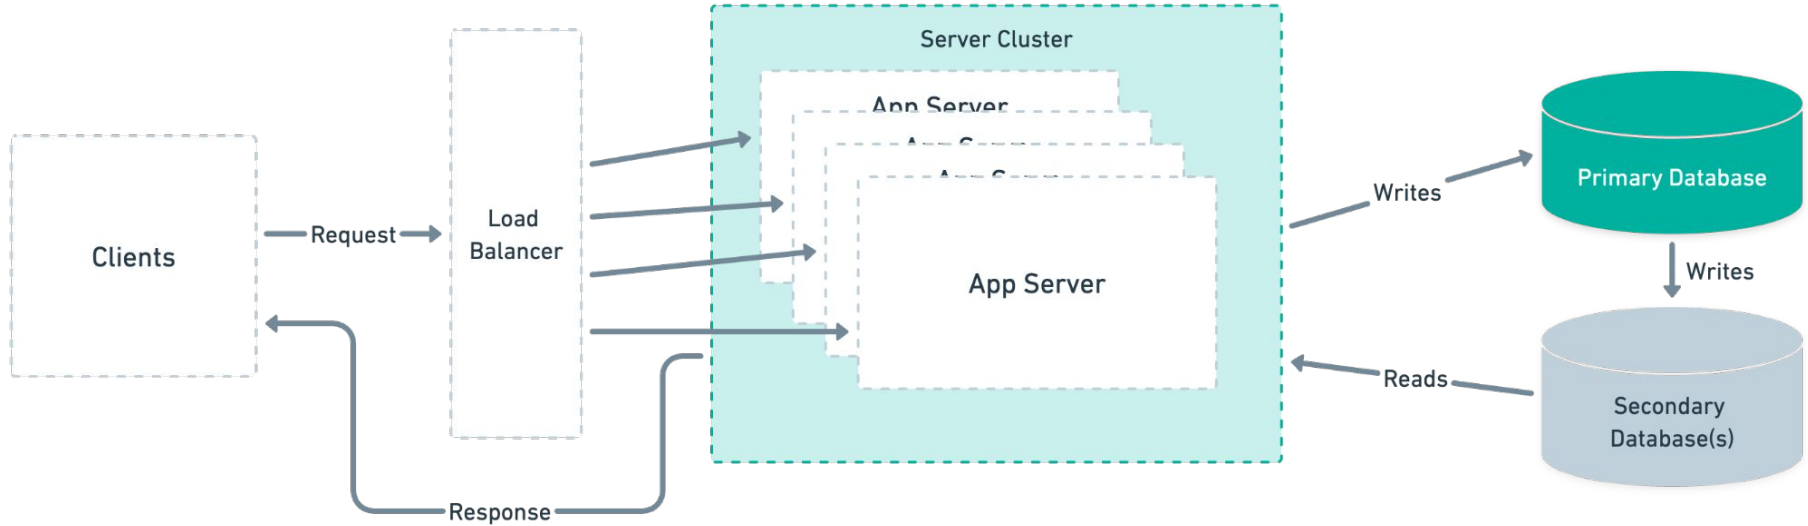
Database Replication



Database Replication improves Availability



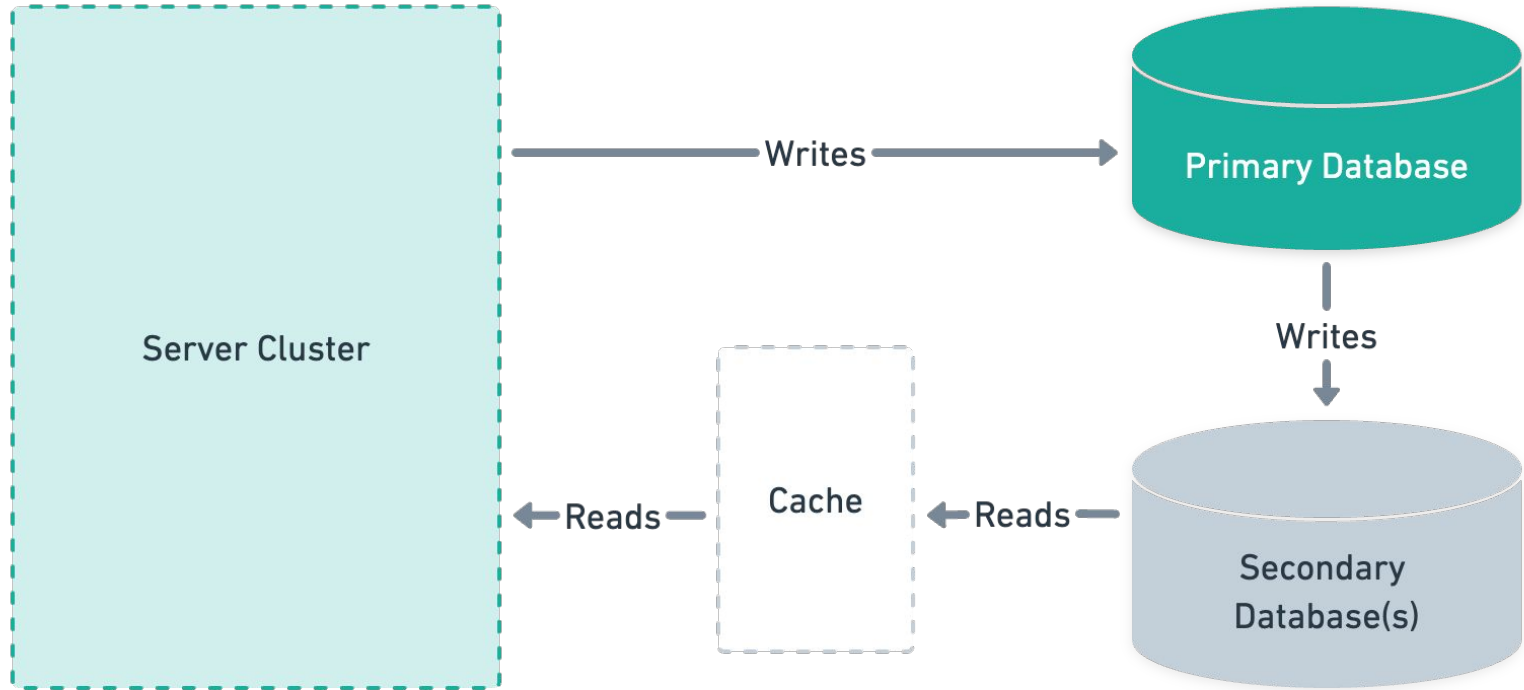
Anything more we can do?



Caching

- Even with replication, calls to the database are expensive
- Let's read from something faster, when we can

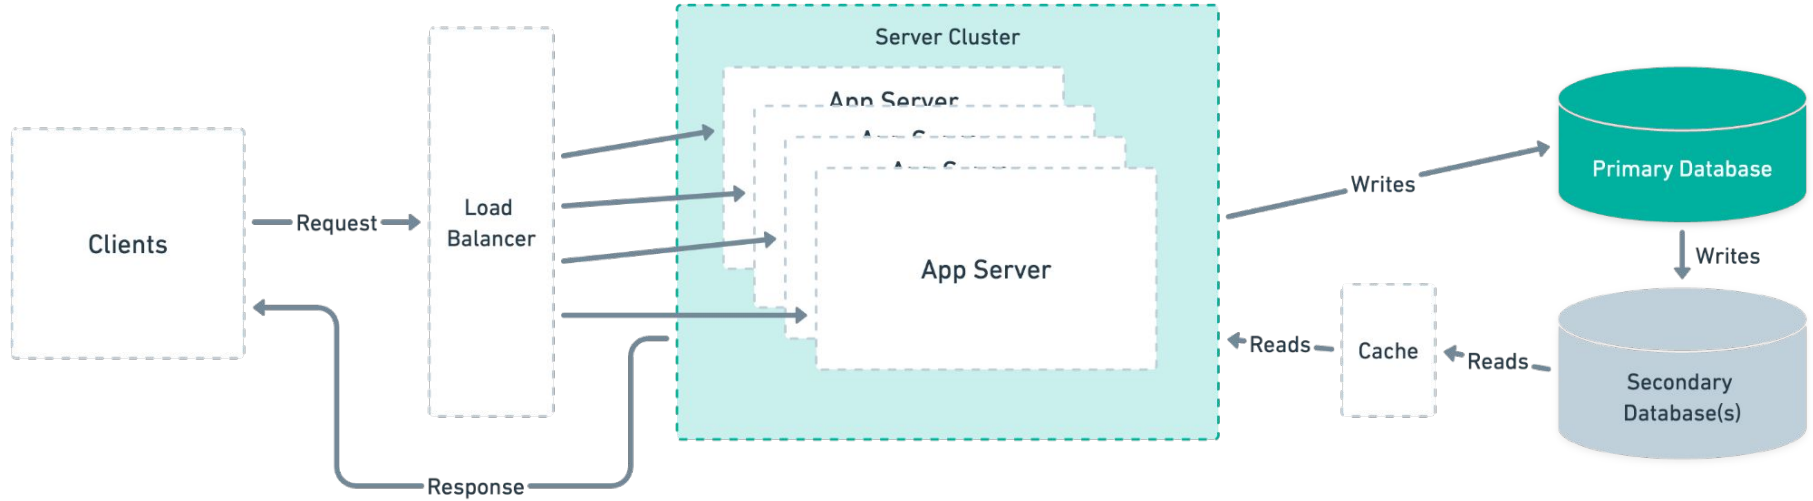
Cache



Why are caches faster than the database?

- Caches are designed to be fast to read from and write to
 - Relational Databases, in particular, require more overhead
- Data is stored in memory (instead of on the disk)
 - Memory is faster than the disk
- They, too, can be horizontally scaled

Anything more we can do?



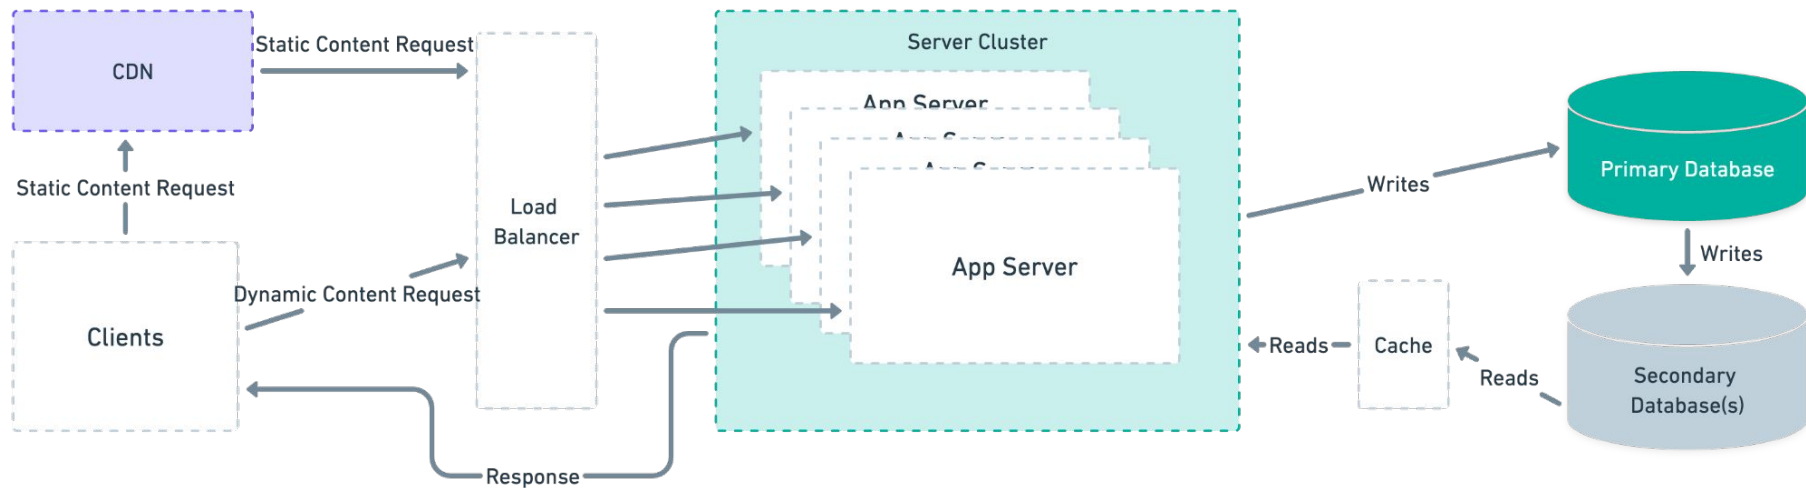
Some Insights

- The information being carried over a network is a physical thing like anything else.
- Imagine if we're passing mail through a pneumatic tube:
 - The further the distance between the sender and receiver, the longer it takes for the mail to arrive.
 - The same is true of electrons or photons over a cable.
- Given that, would storing the data closer to our clients speed up responses?

Content Delivery Network (CDN)

- Network of servers located in different regions of the world
- Direct users to the CDN instance closest to them
- Static content can be cached there
 - Media
 - Images
 - Video
 - Scripts (JS)
 - CSS
- Dynamic content too, in certain cases
- Reduces latency
- Reduces number of requests made to your servers

CDN



How Does Software Talk to Other Software?

- What does it mean, exactly, for a client to make a request to the server?
- What does it mean, exactly, for a server to respond to a request?

Weather Application

- What might the request and response look like for a weather application hosted on the web?
- Request
 - Where do we send this request?
 - <http://www.weather-info.com/weather>
 - How?
 - [HTTP GET method](#)
 - How do we tell the app the location for which we want the weather?
 - <http://www.weather-info.com/weather?zip=11210>
- Response
 - Structured data
 - JSON
 - { "temp": 76, "humidity": 40, "precip": false }

Application Programming Interface (API)

- Set of rules and protocols that allows one software application to interact with another
- How one piece of software can communicate with another
- What's an API you've all had experience using?

Weather API

API Documentation

Endpoint	/weather
HTTP Method	GET
Parameters	<code>integer</code> zip
JSON Response	<code>integer</code> temp <code>string</code> humidity <code>string</code> precip


Request Code

```
public static void main(String[] args) {  
    String url = "/weather/?zip=11210";  
    Response response = Request.Get(url).execute();  
    String output = response  
        .returnContent()  
        .asString();  
    System.out.println(output);  
}
```

Response

```
{  
    "temp": 76,  
    "humidity": 40,  
    "precip": false  
}
```

Where is your API implemented?




REST API

- Clients use [HTTP methods](#) such as GET, PUT, DELETE, etc. to access server data.
- Stateless

Why does Statelessness matter?

- Each request can be treated independently
- A request can be routed to any one of the servers



System Design Interview

<https://www.amazon.com/System-Design-Interview-insiders-Second/dp/B08CMF2CQF>



System Design Interview Framework - Alex Xu

1. Understand the problem & establish design scope
2. Propose high-level design & get buy-in
3. Design deep-dive
4. Wrap-up

Step 1: Understand the Problem & Establish Design Scope

1. Functional requirements

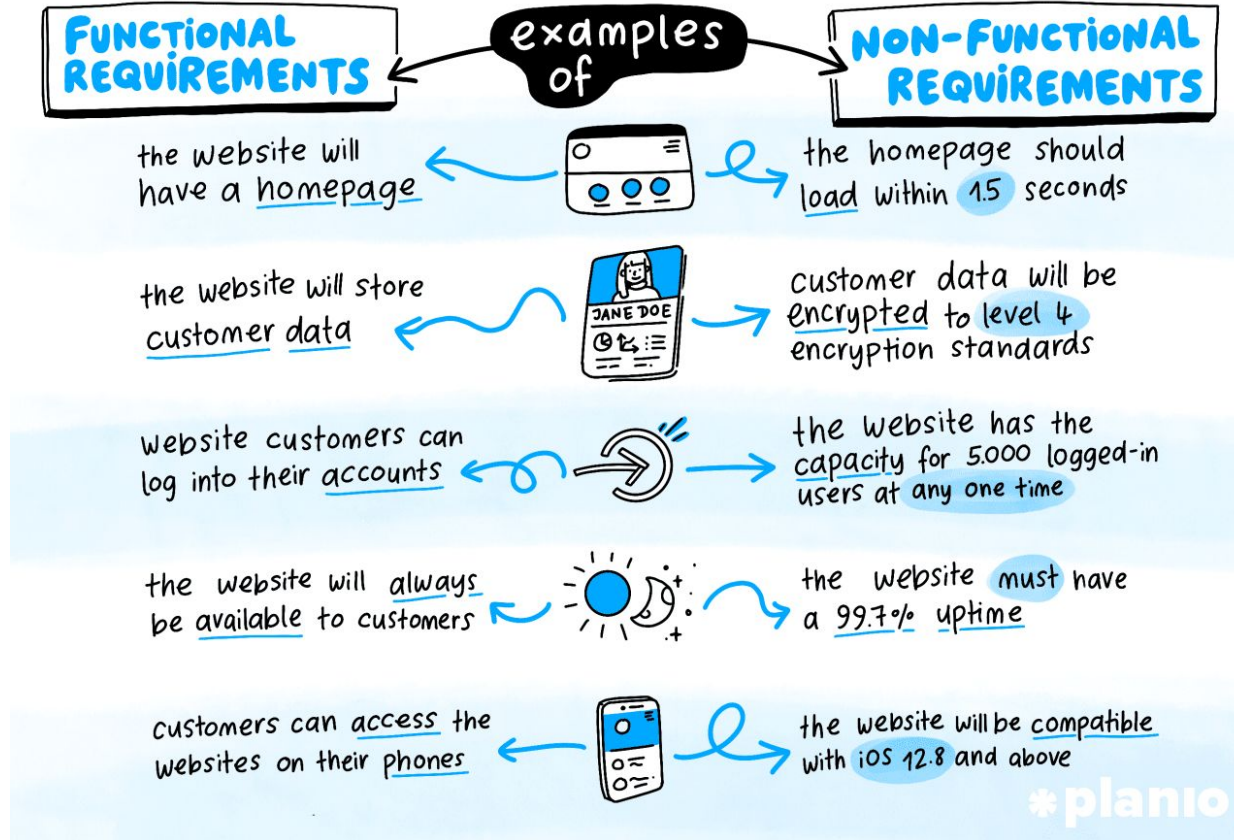
- What the System Should Do?
 - Narrow down the scope as much as possible
 - e.g. If the question is "Design Twitter", specify which exact parts
 - Timeline
 - Followers
 - DMs
 - Search
 - Media
 - Authentication
- Explicitly list use cases
 - Logged In Users
 - Can post a tweet, can send a message, etc.
 - Logged Out Users
 - Can't post a tweet, can't send a message, etc.

Step 1: Understand the Problem & Establish Design Scope

2. Non-functional requirements

- **Performance**: Describes how well the system performs under certain conditions, including response time, throughput, and scalability.
- **Reliability**: Ensures the system operates correctly and reliably over time, including measures like availability, fault tolerance, and recovery.
- **Scalability**: Addresses the system's ability to handle increased load or growth in terms of users, data, or transactions.
- **Availability**: Specifies the percentage of time the system should be operational and accessible.
- **Security**: Outlines the security measures and controls to protect the system from unauthorized access, data breaches, and other security threats.

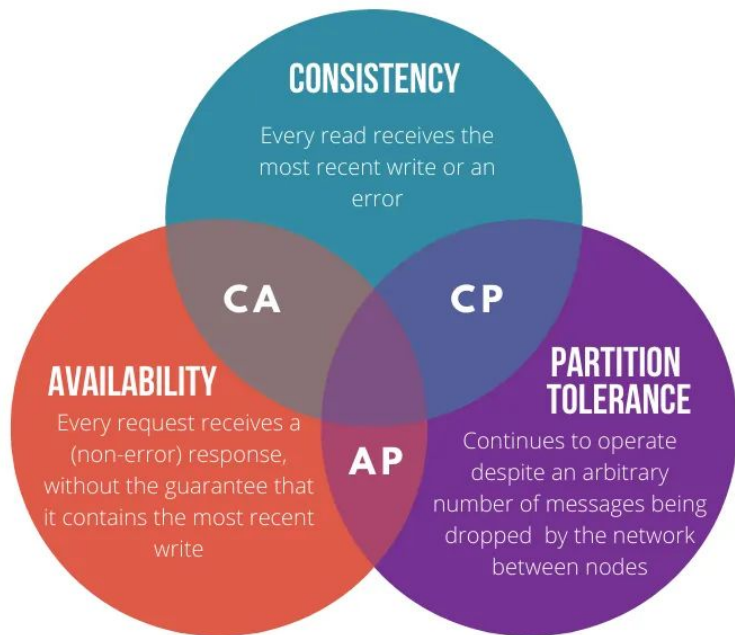
Functional vs Non-Functional Requirements



CAP Theorem

In a **Distributed** system, you can only guarantee two of the following:

- **Consistency**
 - Reads return the exact same data for all users
- **Availability**
 - Every request receives a response
- **Partition Tolerance**
 - Continue working even if two nodes can't communicate



CAP Theorem: Why only 2 out of 3?

For Distributed Systems, Partition Tolerance is considered a necessity. So if a partial network outage occurs during a read/write, your system must continue operating. Do you:

- Cancel the operation?
 - Decreases **availability** but ensures **consistency**.
- Proceed with the operation?
 - Provides **availability** but risks **inconsistency**.

CAP Theorem

During interviews, be explicit about which choice you're making and why.

- Choose **consistency** over availability when
 - Data integrity is a high priority e.g. financial transactions
 - Read-heavy Systems
 - If writes are rare, there will be few disruptions anyway
- Choose **availability** over consistency when
 - Data integrity isn't essential e.g. Two users seeing a different number of likes on a tweet
 - User experience would be greatly harmed by unavailability
 - Or not harmed by temporary inconsistency (a.k.a. Eventual Consistency)
 - Write-heavy Systems
 - Writes are common, so disruptions would be common if consistency were chosen over availability

Step 1: Understand the Problem & Establish Design Scope

3. Back-of-the-Envelope Estimations


- Estimates you create to get a good feel for which designs will meet your requirements
 - Use a combination of thought experiments and common performance numbers
 - Pick whole numbers. Accuracy isn't important, you're only worried about the order of magnitude.
- Load
 - Requests Per Second
 - Data Volume
 - User Traffic
- Storage
 - Amount of Storage Required
- Resources:
 - Number of:
 - Servers
 - CPUs
 - Memory
- Network Bandwidth
- Latency

Step 1: Back-of-the-Envelope Estimations for Twitter

- Assumptions:
 - 300 million monthly active users.
 - 50% of users use Twitter daily.
 - Users post 2 tweets per day, on average.
 - 10% of tweets contain media.
 - Data is stored for 5 years.
- Estimations:
 - Query per second (QPS) estimate:
 - Daily active users (DAU) = $300 \text{ million} * 50\% = 150 \text{ million}$
 - Tweets QPS = $150 \text{ million} * 2 \text{ tweets} / 24 \text{ hour} / 3600 \text{ seconds} = \sim 3500$
 - Peek QPS = $2 * \text{QPS} = \sim 7000$
 - Media Storage estimate:
 - Average tweet size:
 - tweet_id = 64 bytes
 - text = 140 bytes
 - media = 1 MB
 - Media storage: $150 \text{ million} * 2 * 10\% * 1 \text{ MB} = 30 \text{ TB per day}$
 - 5-year media storage: $30 \text{ TB} * 365 * 5 = \sim 55 \text{ PB}$

Step 2: Propose High-Level Design & Get Buy-In

1. Sketch a simple diagram of your system's core functionalities
 - Use the requirements you gathered in step one as a checklist



Step 2: Propose High-Level Design & Get Buy-In

2. Define the API

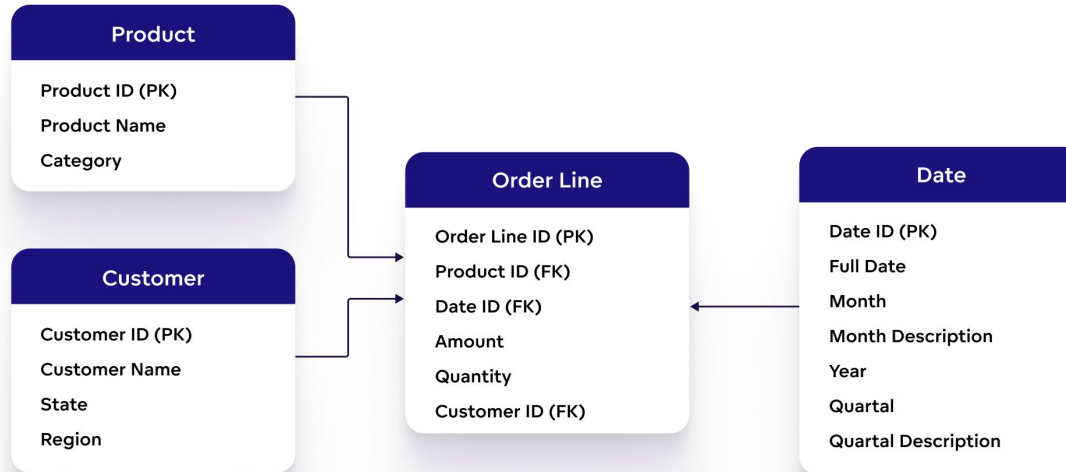
- What endpoints will you have to create?

Endpoint	/weather
HTTP Method	GET
Parameters	<code>integer</code> zip
JSON Response	<code>integer</code> temp <code>string</code> humidity <code>string</code> precip

Step 2: Propose High-Level Design & Get Buy-In

3. Define the Data Model

- What tables will you need to create in your database?
- What are the relationships between those tables?



Step 2: Propose High-Level Design & Get Buy-In

4. Get Buy-In

- Check in with your interviewer, confirm you're on the right track
- Make space for them to interject

Step 3: Design Deep-Dive

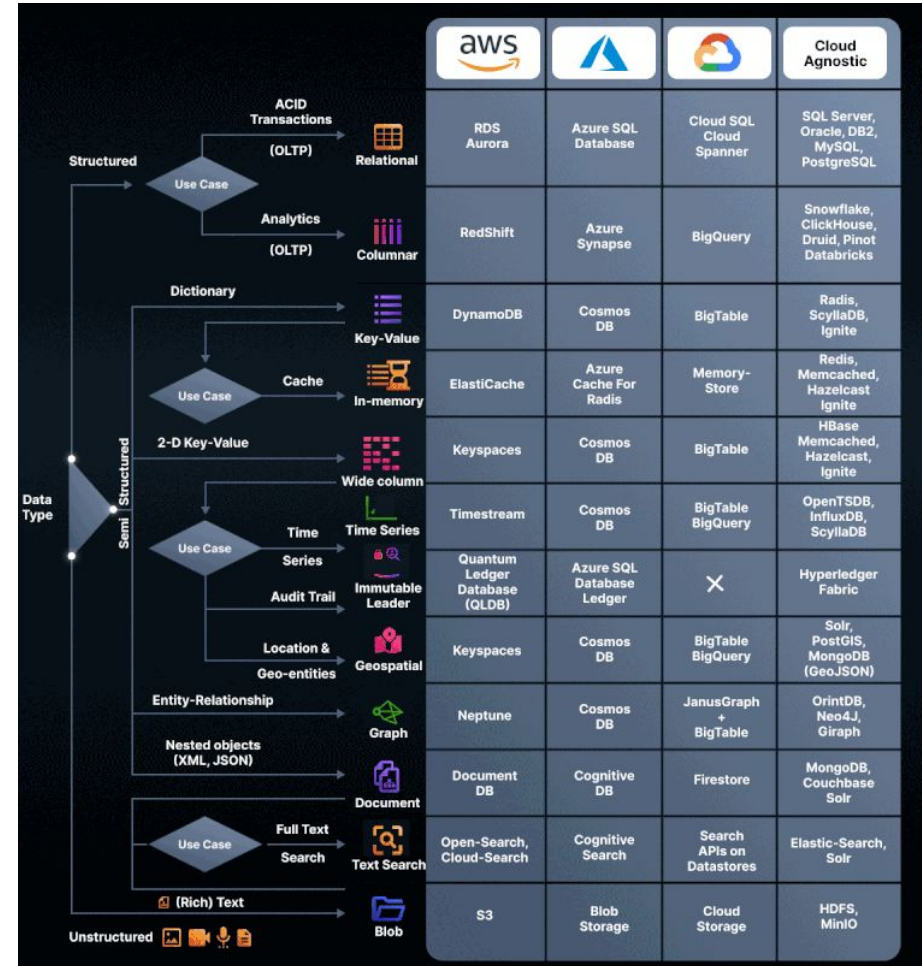
- Make the system:
 - Faster
 - CDNs, Caches
 - Robust
 - Database Replication
 - Sharding
 - Secure
 - OAuth
 - ACLs
 - Encryption
- } Also makes it faster

Step 3: Design Deep-Dive

- Technology Choices
 - What sort of database?
 - Relational vs NoSQL

Choice of Database System

- Choose an appropriate database technology based on the type of data you're dealing with.
- Relational is often a safe bet
- Blob storage for media
- NoSQL for performance / non-relational data



Step Four: Wrap-Up

- Give the interviewer a recap of your design
- Error Cases
- Edge Cases
- Metrics, Monitoring, Alerting
- Make space for the interviewer to pick something to drill into