

# SYSTEMS DESIGN

CARLOS CUEVAS

ADJUNCT LECTURER - BROOKLYN COLLEGE

# Typical Software Engineer Interview

Typically you'll be given several distinct rounds of interviews:

- Data Structures & Algorithms
  - [Leetcode](#) style questions
- Behavioral
  - "Tell me about a time you disagreed with someone..."
- Systems Design
  - Primarily asked of Senior Engineers but it's become increasingly common to ask even entry level engineers

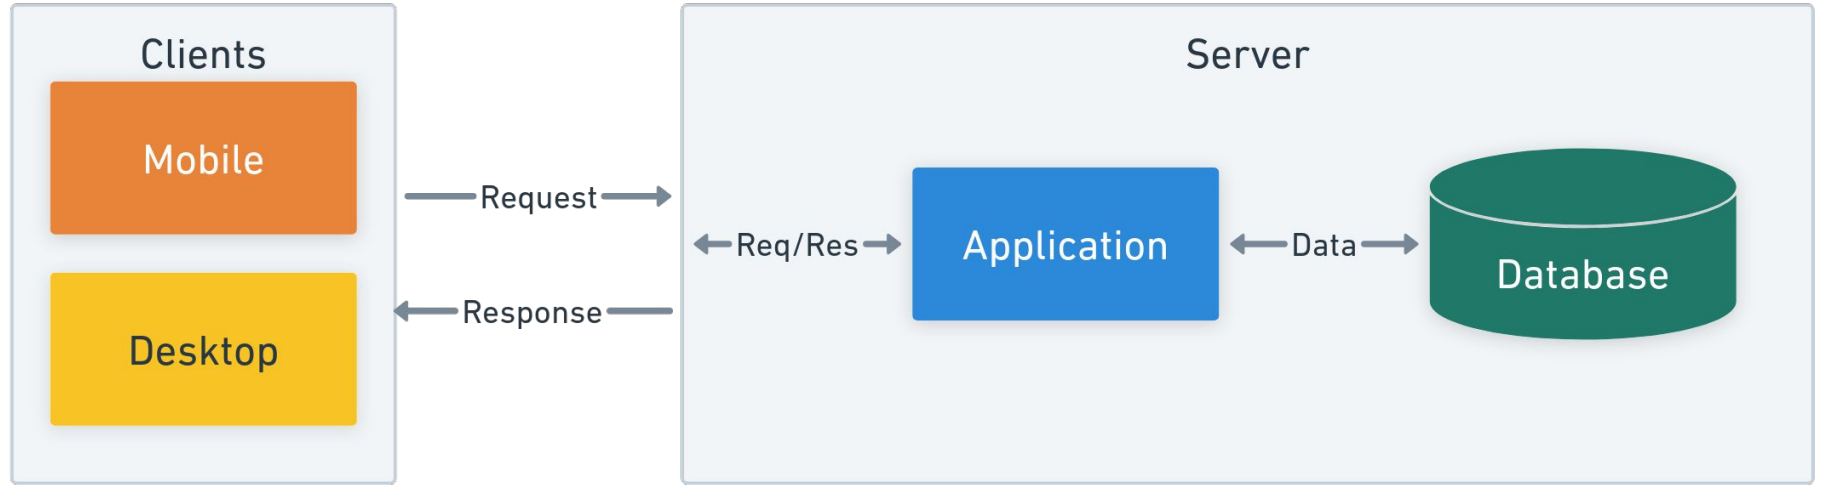
# Systems Design Interview

- You'll be given a purposefully broad prompt.
- Design a system that accomplishes x. For example:
  - Chat System
  - URL Shortener
  - Youtube
  - Many Others
- It's on you to narrow scope as much as possible, but we'll get to that later.

# Building Blocks

- Many systems have overlap
- Let's learn some common components of many large scale systems

# Client - Server Model



# How Does Software Talk to Other Software?

- What does it mean, exactly, for a client to make a request to the server?
- What does it mean, exactly, for a server to respond to a request?

# Weather Application

- What might the request and response look like for a weather application hosted on the web?
- Request
  - Where do we send this request?
    - <http://www.weather-info.com/weather>
  - How?
    - [HTTP GET method](#)
  - How do we tell the app the location for which we want the weather?
    - <http://www.weather-info.com/weather?zip=11210>
- Response
  - Formatted data
    - JSON
      - { "temp": 76, "humidity": 40, "precip": false }

# Application Programming Interface (API)

- Set of rules and protocols that allows one software application to interact with another
- i.e. How one piece of software can communicate with another
- What's an API you've all had experience using?



# Weather API

## API Documentation

Endpoint	/weather
HTTP Method	GET
Parameters	<code>integer</code> zip
JSON Response	<code>integer</code> temp <code>string</code> humidity <code>string</code> precip

## Request Code

```
public static void main(String[] args) {  
    String url = "/weather/?zip=11210";  
    Response response = Request.Get(url).execute();  
    String output = response  
        .returnContent()  
        .asString();  
    System.out.println(output);  
}
```

## Response

```
{  
    "temp": 76,  
    "humidity": 40,  
    "precip": false  
}
```

# Weather API

## Response Code

```
public static void respond(Request request) {  
    String zip = getQueryParam("zip", request);  
  
    String weatherData = getWeatherData(zip);  
  
    Response response =  
        Response.ok(weatherData, MediaType.APPLICATION_JSON).build();  
  
    response.returnContent().writeTo(System.out);  
}
```

# What exactly is a Server?

- The term is overloaded
- Hardware
  - A physical\* computer
  - on which **server software** is running
- Software
  - A process that is listening for requests and responding to those requests (i.e. serving)

\* Could be:

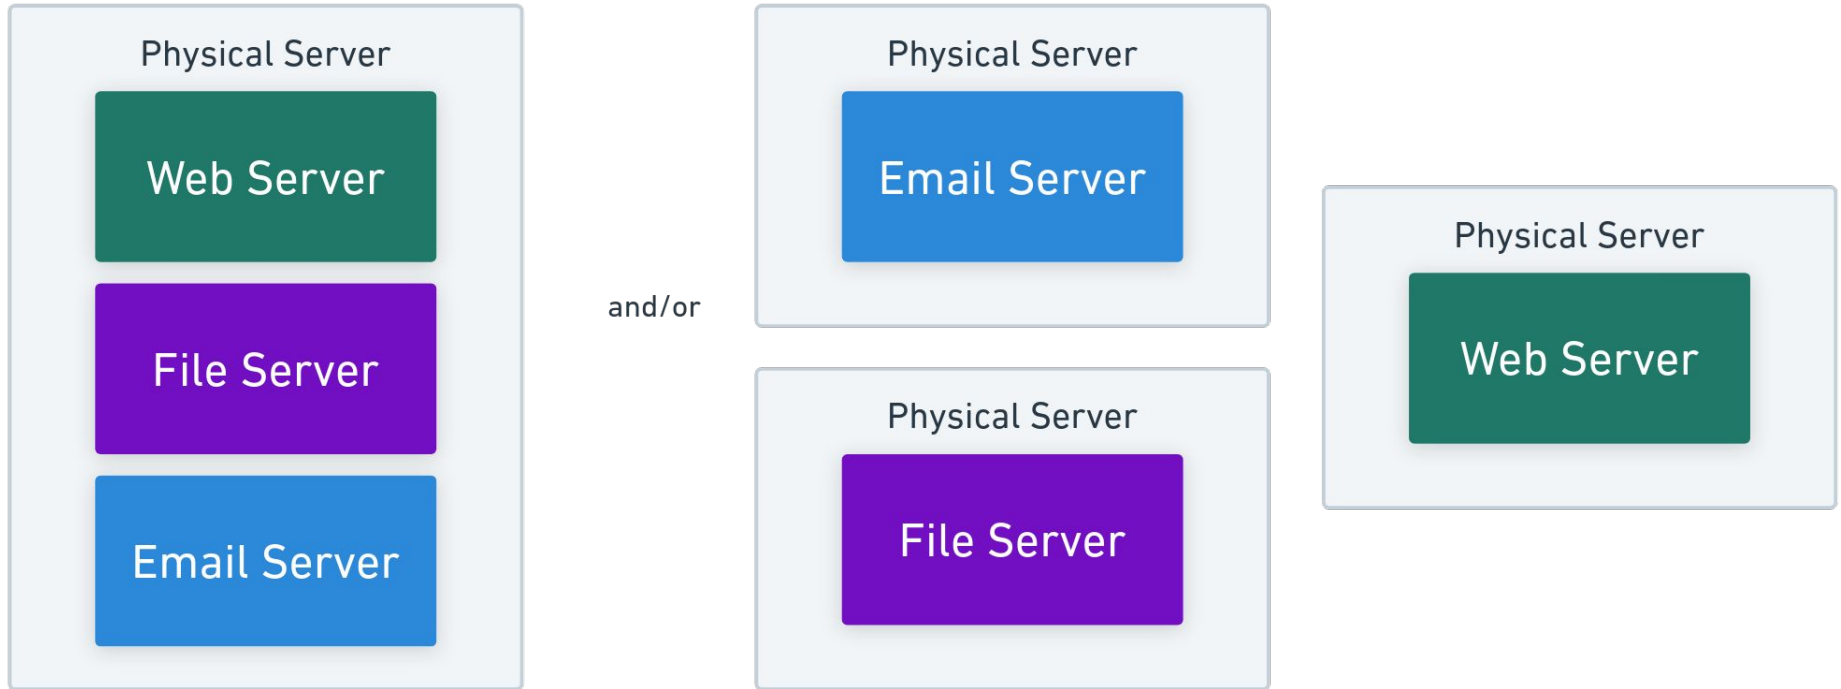
- ❖ Physical Computer
- ❖ Virtual Machine / Container
- ❖ Cloud Service

# Server Software

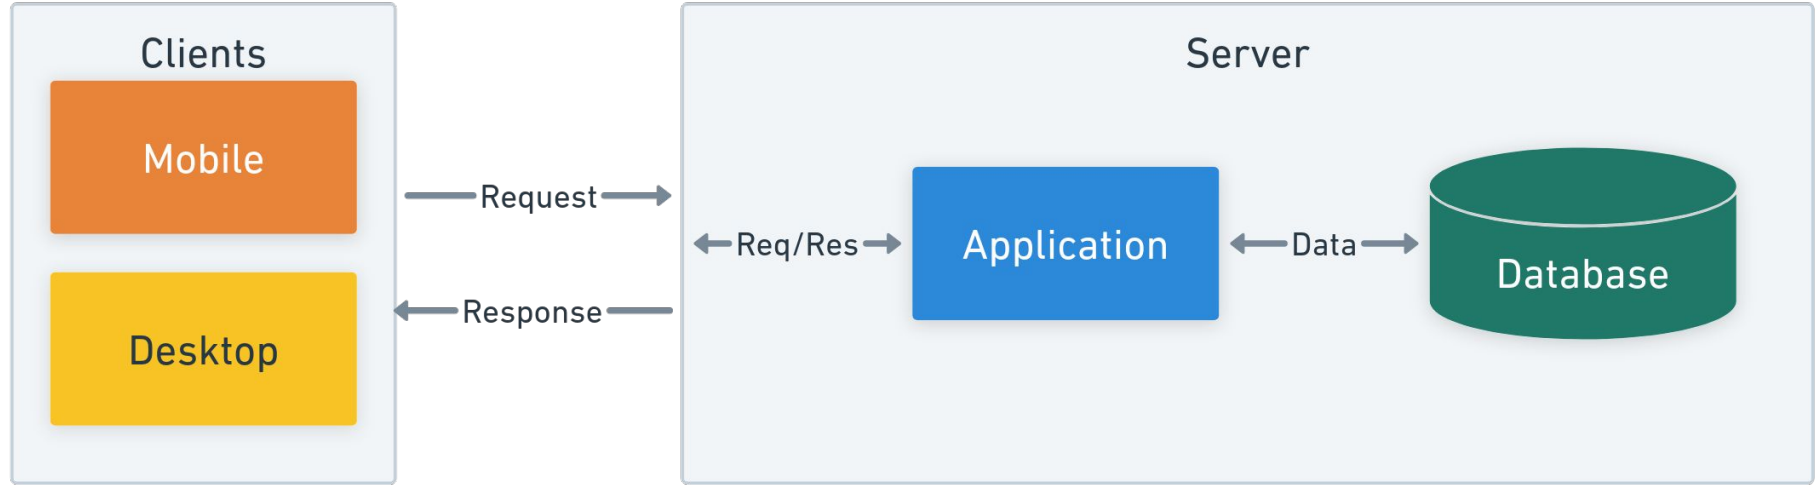
- Web / HTTP
  - Apache, Nginx
- File
  - FTP, Samba
- Email
  - Microsoft Exchange
- Database
  - MySQL, Oracle
- Many, many more

# Hardware & Software


- One Physical Server can run multiple server applications
- Or, you can run multiple physical servers each running its own application.



# Single Server





# Multi-Server Setup



# Horizontal vs Vertical Scaling

- Vertical: a more powerful computer
- Horizontal: more computers





# Horizontal vs Vertical Scaling


- There's a limit on how much you can vertically scale
  - After you've maxed out the most powerful CPU in existence, then what?
- Horizontal Scaling is, essentially, infinite
  - Just add more instances and distribute the load

# Parallelization


- You have 10 washing machine size bags of laundry to do
- With 1 washing machine, you can do 1 at a time
- With 10 washing machines, you can do 10 at a time (in parallel)
- Horizontal Scaling: add lots of washing machines

# Virtualization & Containerization

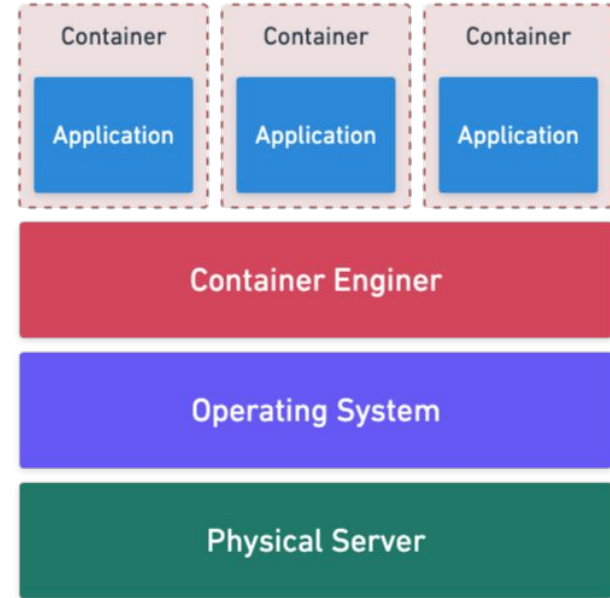
## Bare Metal



## Virtualization



## Containerization



Which requires more knowledge of what's under the hood?



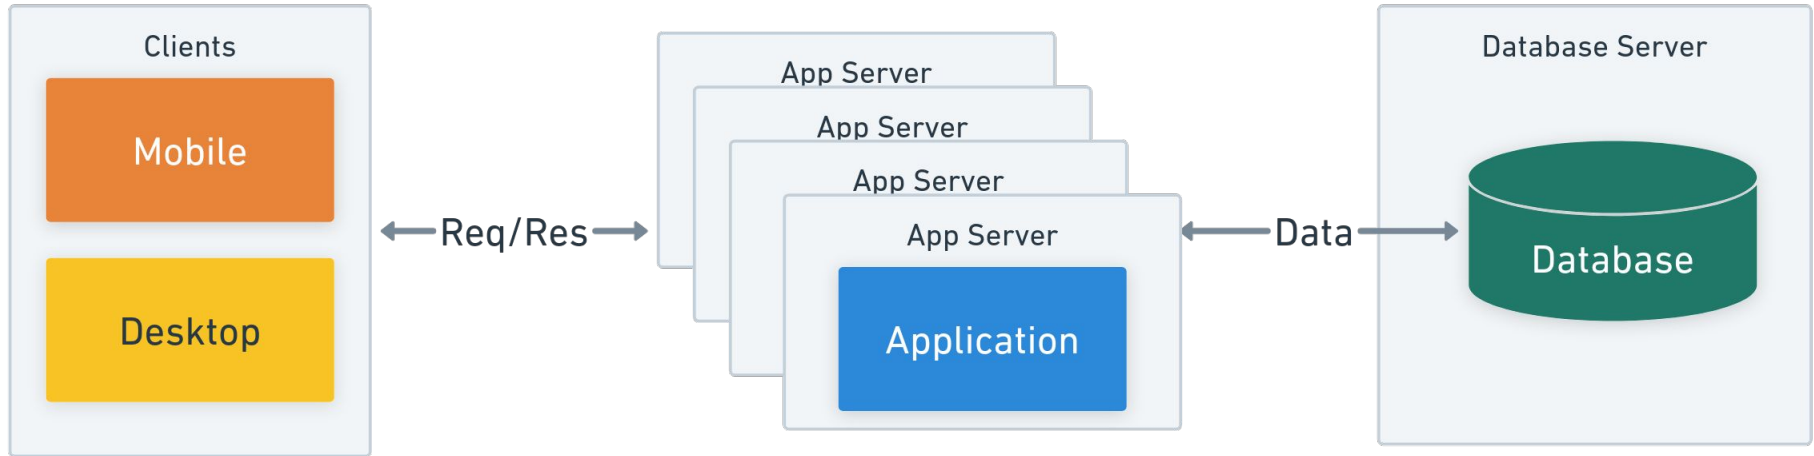
# Abstraction

- Hiding details
- Why bother?
  - Increases ease of use / reduces complexity (for the user, at least)
  - Increases portability / looser coupling
- Trade offs
  - Limited control
  - Increased overhead for the one implementing the abstraction layer

# Containerized Applications

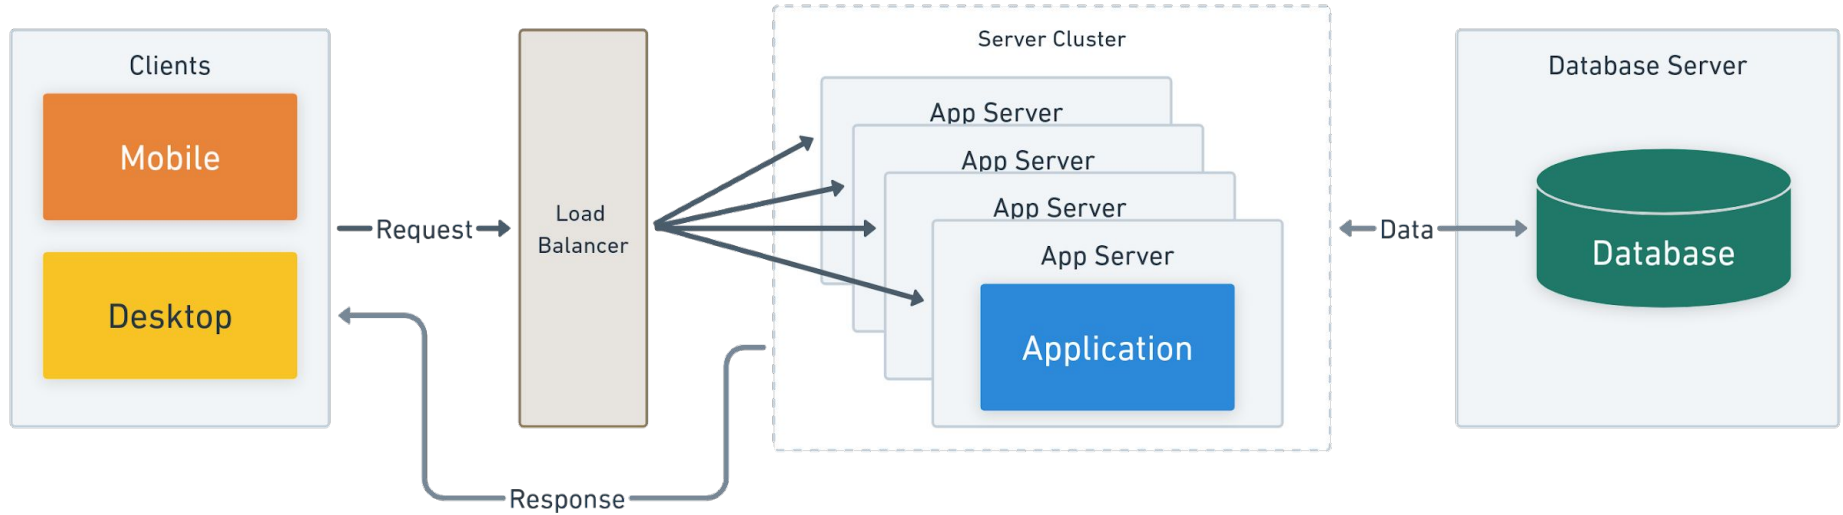
- Containers can be easily spun up and down in milliseconds, enabling you to quickly react to fluctuating demands
- Containers are self-contained units that bundle an application with all its dependencies. This allows them to run consistently across different environments, regardless of the underlying infrastructure.

# Horizontal Scaling



I've added more servers, but now what? How do we make sure the same server isn't used for every request?

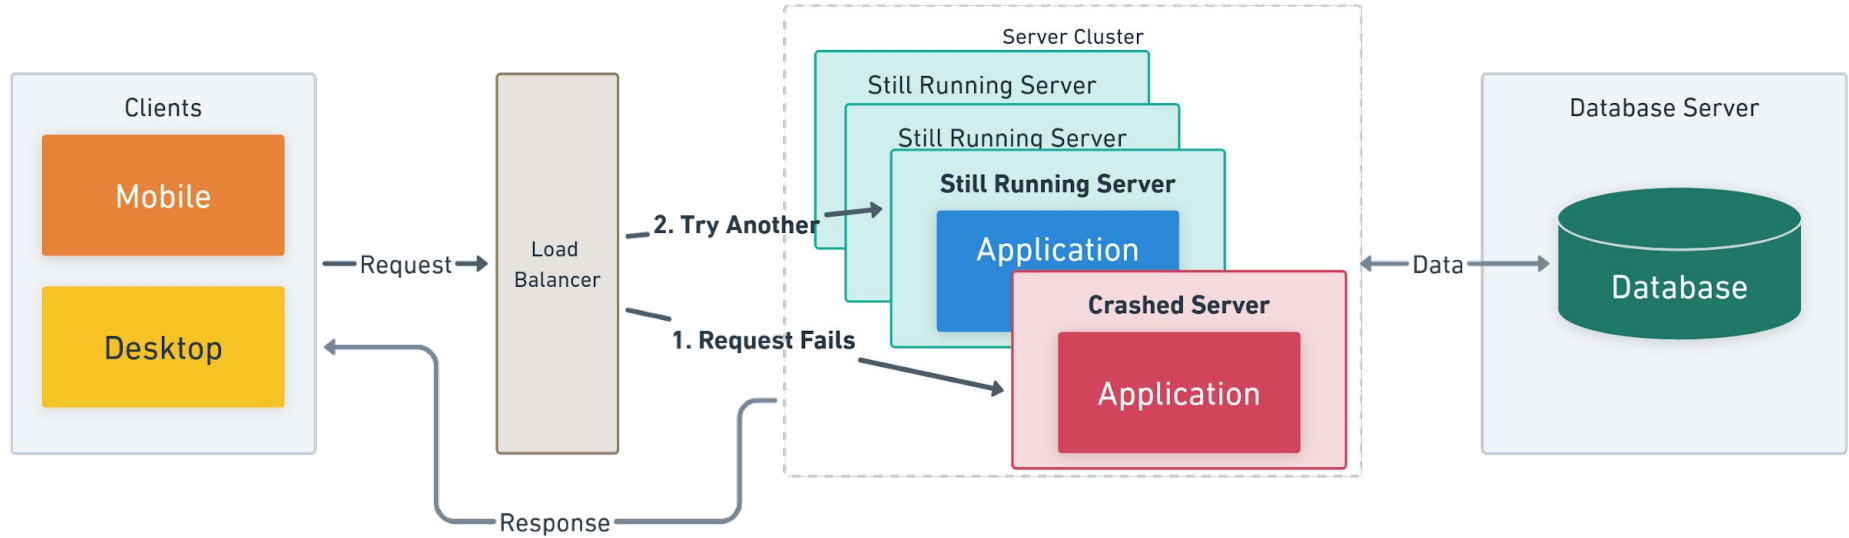
# Load Balancing



Load Balancers ensure an even distribution of traffic.



# Horizontal Scaling + Load Balancing = Availability




# REST API

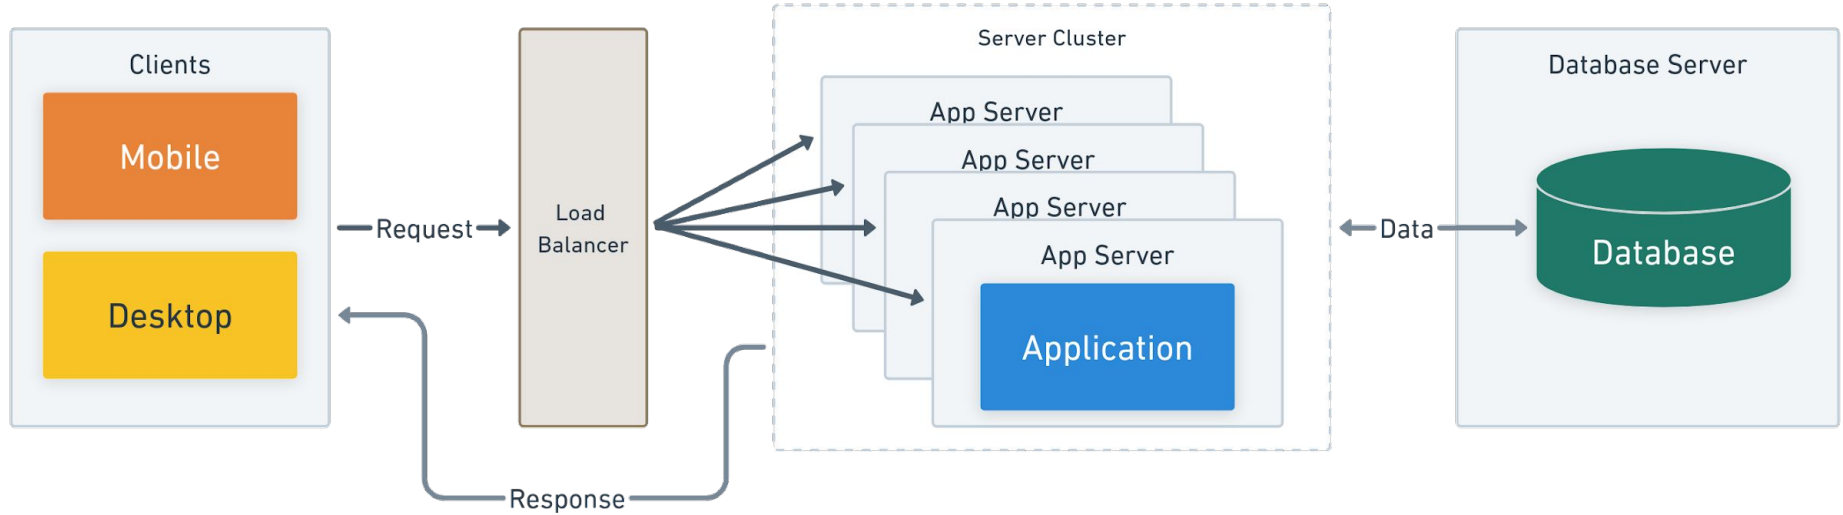
- A style of API design
- Clients use [HTTP methods](#) such as GET, PUT, DELETE, etc. to access server data.
- Stateless
  - Each request is processed purely based on the information provided within the request
  - The server doesn't "remember" anything about a client's previous requests

# Why does Statelessness matter?


- Since each request can be treated independently, a request can be routed to any one of the servers



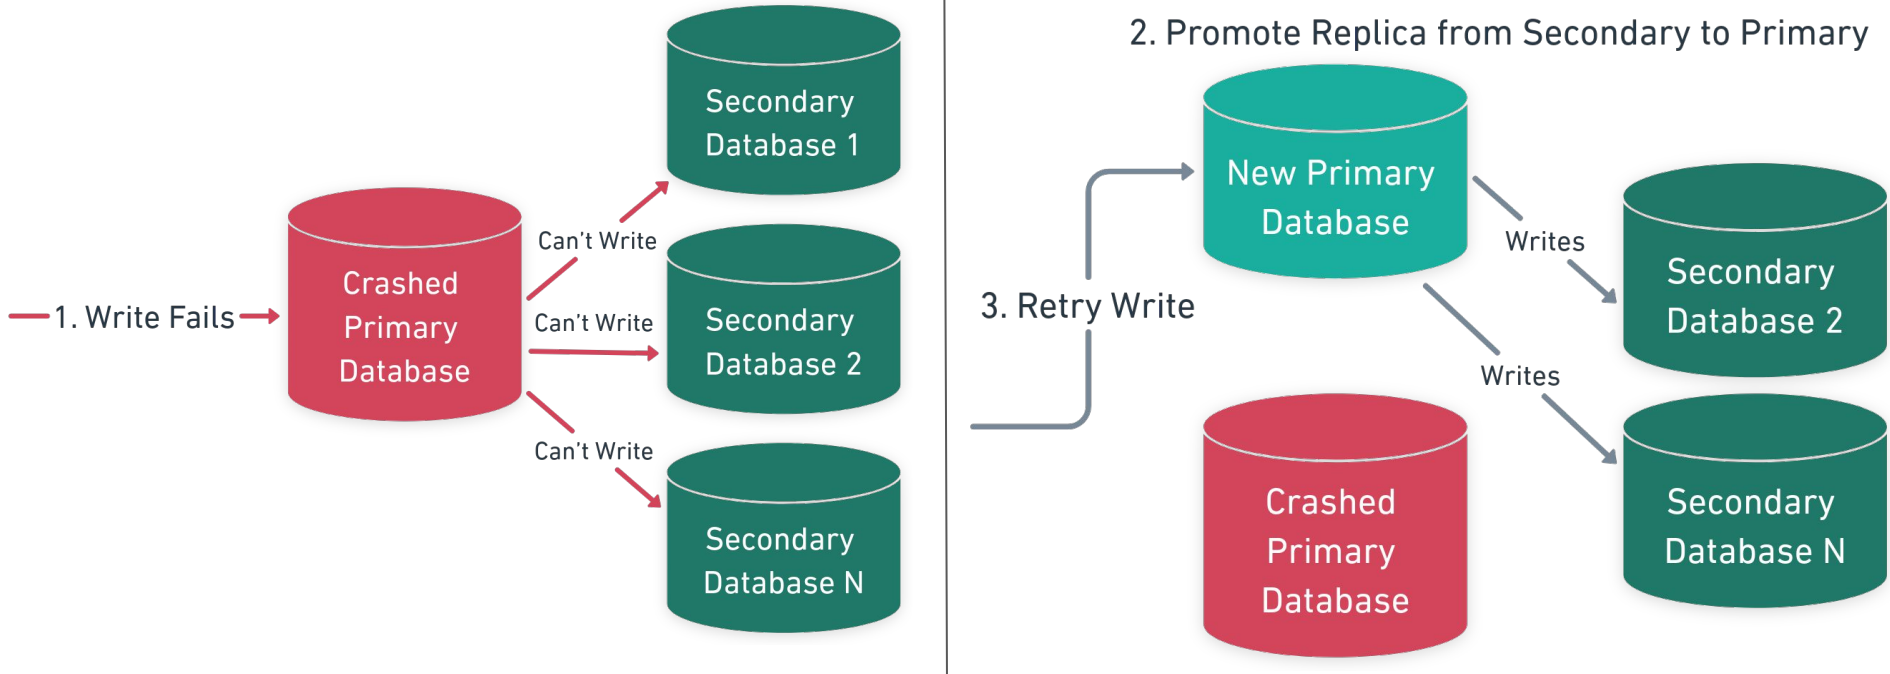
# What else can we scale up?




# Database Replication



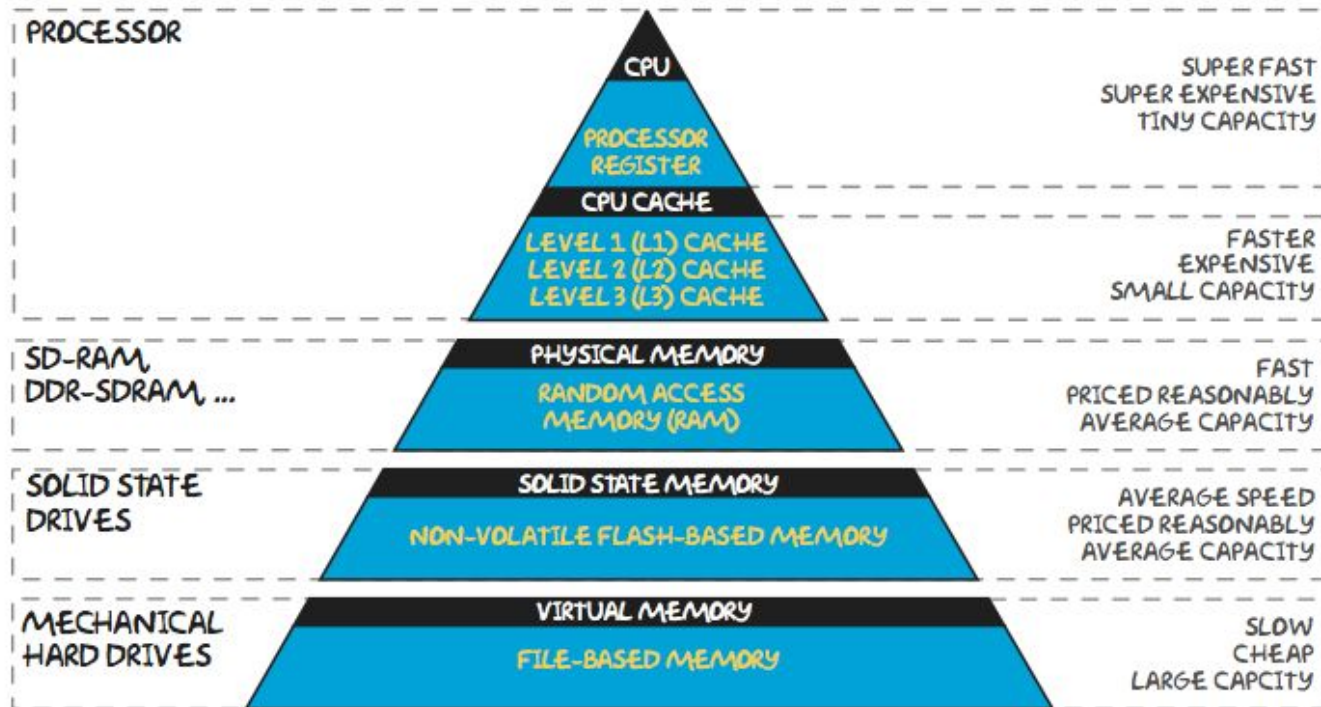
# Database Replication improves Availability



# Anything more we can do?



# Databases are fast, but we can do better

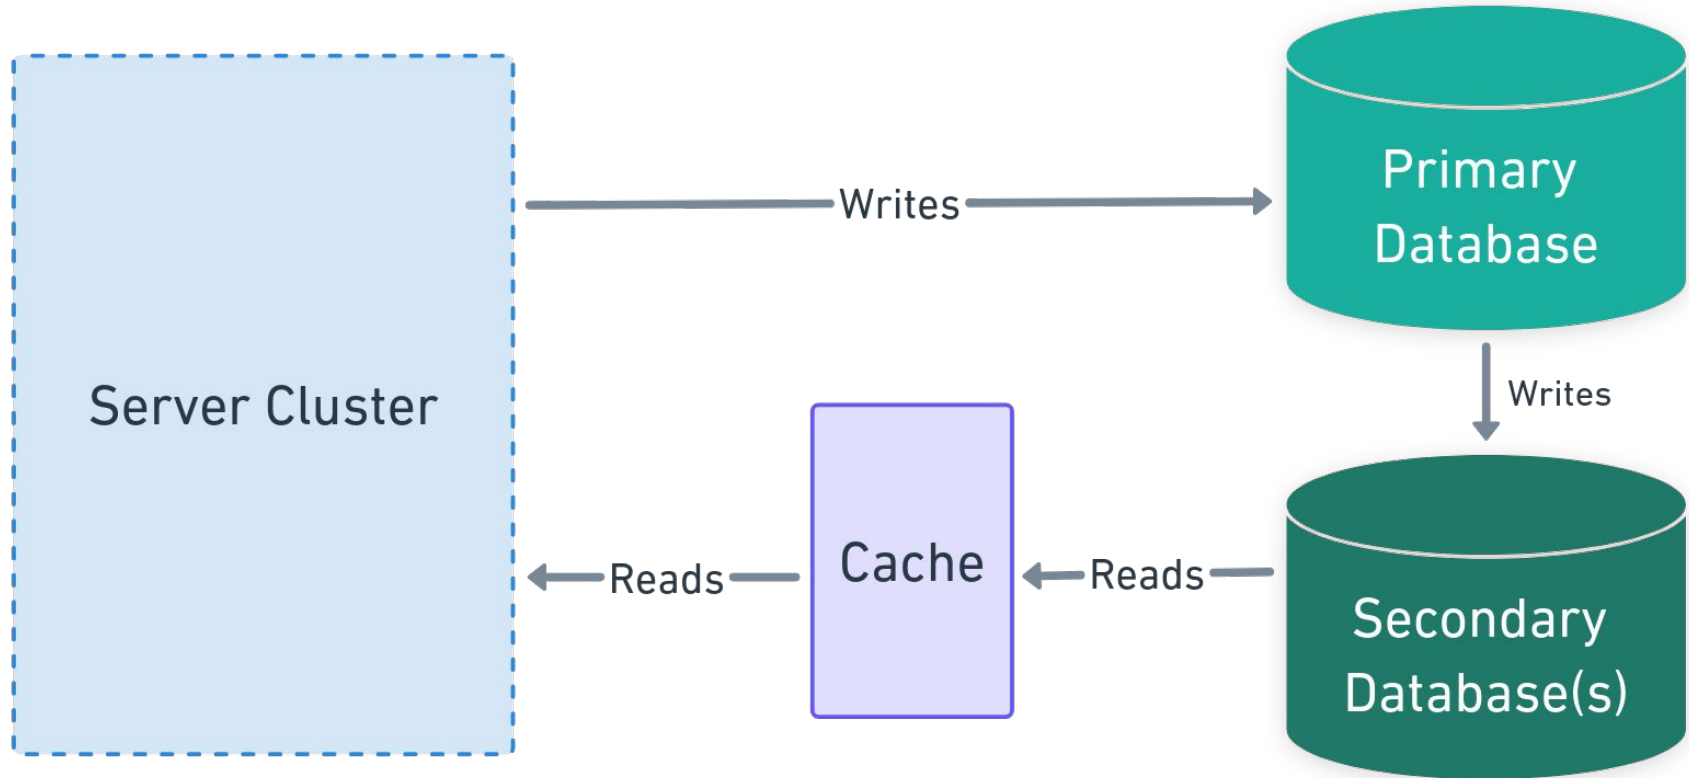




# Caching

- Even with replication, calls to the database are expensive
- Let's read from something faster, when we can


# Cache



# Why are caches faster than the database?


- Caches are designed to be fast to read from and write to
  - Caches tend to be simple key/value stores.
  - Relational Databases enforce constraints and support sophisticated relationships between sets of data.
- Data is stored in memory (instead of on the disk)
  - Memory is faster than the disk.
- They, too, can be horizontally scaled.

# Anything more we can do?



# Some Insights


- The information being carried over a network is a physical thing like anything else.
- Imagine if we were passing mail through a pneumatic tube:
  - The further the distance between the sender and receiver, the longer it takes for the mail to arrive.
  - The same is true of electrons or photons over a cable.
- Given that, would storing the data closer to our clients speed up responses?



# Content Delivery Network (CDN)

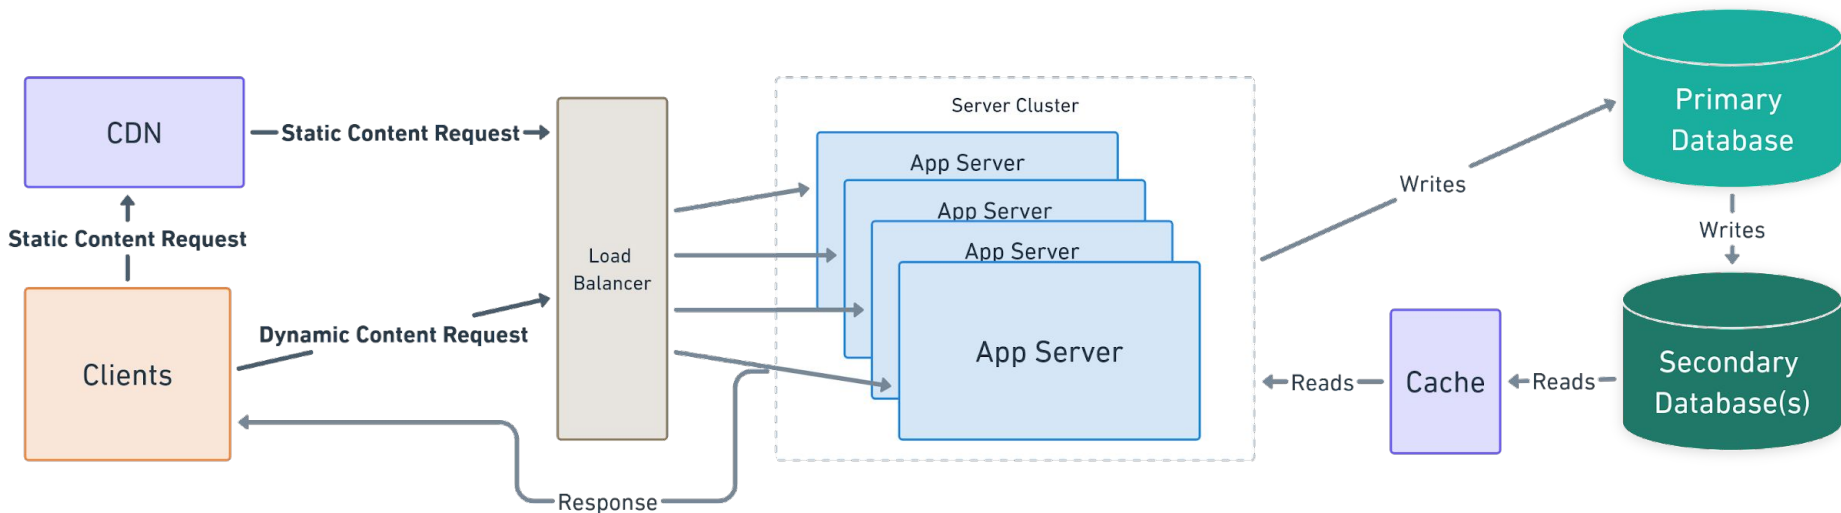
- Networks of servers located in different regions around the world
- Direct users to the CDN instance closest to them
- Static content can be cached there
  - Media
    - Images
    - Video
  - Scripts (JS)
  - CSS
- Dynamic content too, in certain cases
  - If it changes infrequently
- Reduces latency
- Reduces number of requests made to your servers

# CDN



## Anything more we can do?

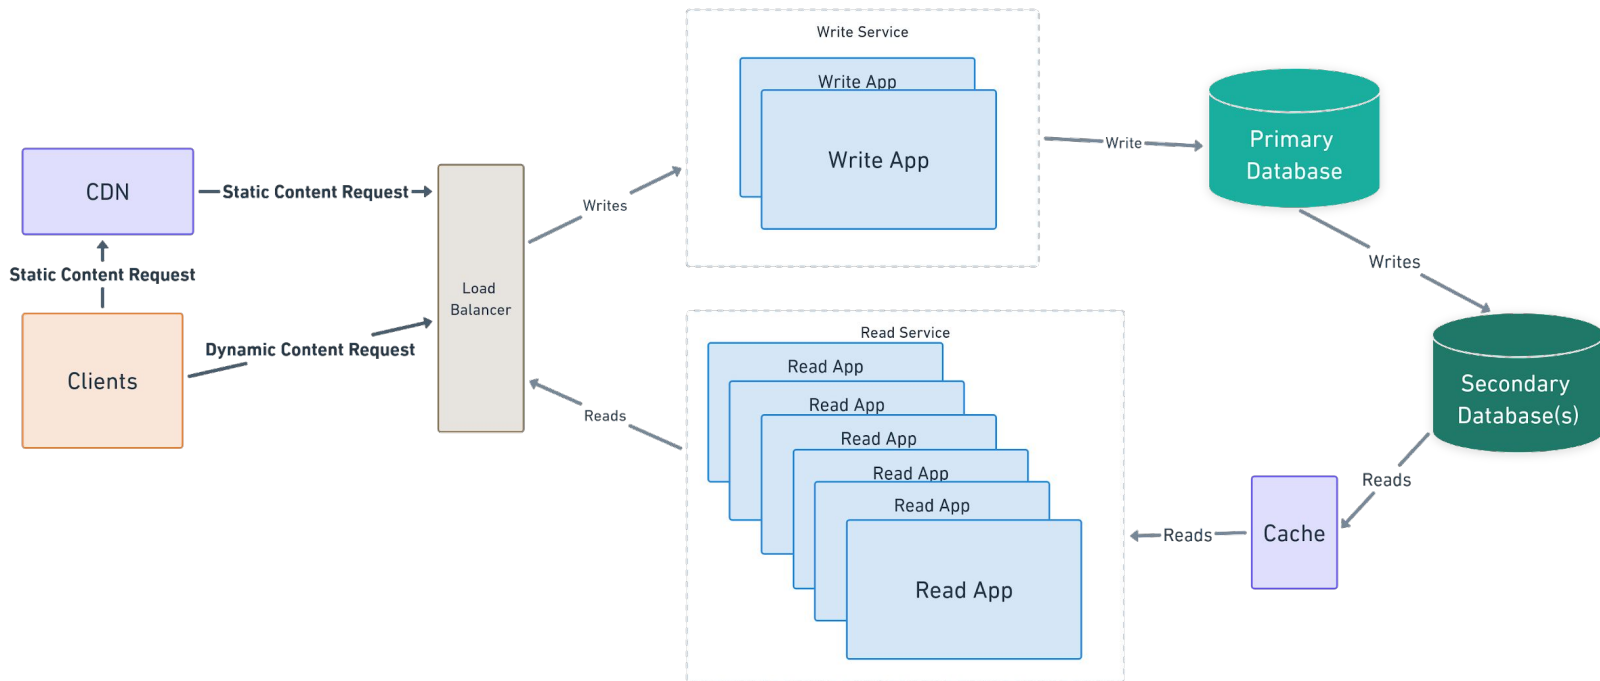
- What if our application does more than one thing, but a lot more of one out of all the things it does? (e.g. way more reads than writes or vice versa)





# Services

- We can break up our application into smaller functions/domains called services and scale those independently as well.



# Systems Design Interview

- It's about the journey, not the destination.
  - There isn't one right answer.
  - Communication is as (if not more) important than your particular design.
  - Particularly, demonstrating an understanding of your decisions, the tradeoffs, etc.
- Above all, do not jump straight into designing. Ask **a lot** of questions first.

# System Design Interview

<https://www.amazon.com/System-Design-Interview-insiders-Second/dp/B08CMF2CQF>



# System Design Interview Framework - Alex Xu

1. Understand the problem & establish design scope
2. Propose high-level design & get buy-in
3. Design deep-dive
4. Wrap-up

# Step 1: Understand the Problem & Establish Design Scope

## 1. Define the **Functional Requirements**

- **What** the should the system do?
  - Narrow down the scope as much as possible
  - e.g. If the question is "Design Twitter", specify which exact parts
    - Timeline
    - Followers
    - DMs
    - Search
    - Media
    - Authentication
- Explicitly list use cases
  - Logged In Users
    - Can post a tweet, can send a message, etc.
  - Logged Out Users
    - Can't post a tweet, can't send a message, etc.

# Step 1: Understand the Problem & Establish Design Scope

## 2. Define **Non-Functional Requirements**

- **How** the system should do it.
  - **Performance:** Describes how well the system performs under certain conditions, including response time, throughput, and scalability.
  - **Reliability:** Ensures the system operates correctly and reliably over time, including measures like availability, fault tolerance, and recovery.
  - **Scalability:** Addresses the system's ability to handle increased load or growth in terms of users, data, or transactions.
  - **Availability:** Specifies the percentage of time the system should be operational and accessible.
  - **Security:** Outlines the security measures and controls to protect the system from unauthorized access, data breaches, and other security threats.

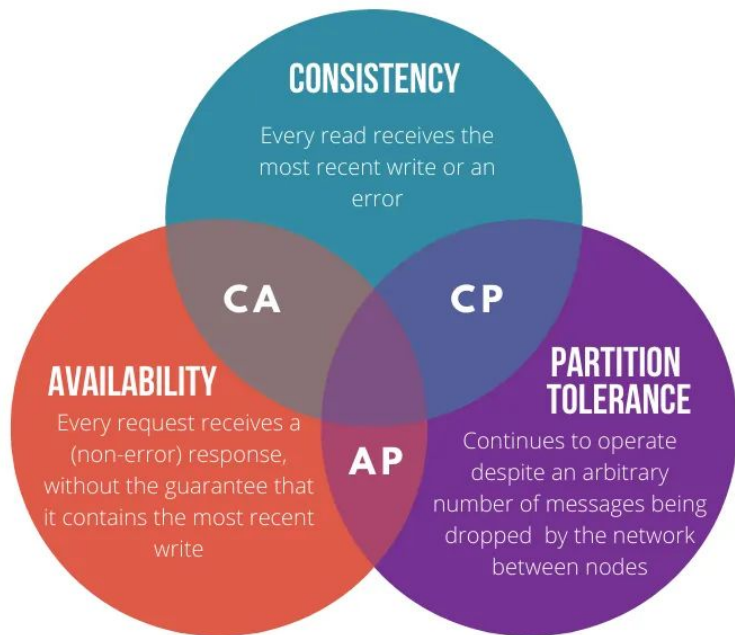
# Functional vs Non-Functional Requirements

Functional Requirements	Non-Functional Requirements	
Users can post tweets	<u>Availability</u>	Highly Available
Users can delete tweets	<u>Latency</u>	$\leq 200$ ms
Users can search tweets	<u>Consistency</u>	Eventual Consistency
Users can view another's timeline	<u>Scalability</u>	Highly scalable

# CAP Theorem

In a **Distributed** system, you can only guarantee two of the following:

- **Consistency**
  - Reads return the exact same data for all users
- **Availability**
  - Every request receives a response
- **Partition Tolerance**
  - Continue working even if two nodes can't communicate





# CAP Theorem: Why only 2 out of 3?

For Distributed Systems, Partition Tolerance is considered a necessity. So if a partial network outage occurs during a read/write, your system must continue operating. Do you:

- Cancel the operation?
  - Decreases **availability** but ensures **consistency**.
- Proceed with the operation?
  - Provides **availability** but risks **inconsistency**.

# CAP Theorem

During interviews, be explicit about which choice you're making and why.

- Choose **consistency** over availability when
  - Data integrity is a high priority
    - Financial transactions
  - Read-heavy systems
    - If writes are rare, there will be few disruptions during a partial network outage
- Choose **availability** over consistency when
  - Data integrity isn't essential
    - Two users seeing a different number of likes on a tweet
  - User experience would be greatly harmed by unavailability
  - Or not harmed by temporary inconsistency (a.k.a. Eventual Consistency)
  - Write-heavy Systems
    - If writes are frequent, there will be many disruptions during a partial outage

# Step 1: Understand the Problem & Establish Design Scope

## 3. Back-of-the-Envelope Estimations

- Estimates you create to get a good feel for which designs will meet your requirements
  - Use a combination of thought experiments and common performance numbers
  - Pick whole numbers. Accuracy isn't important, you're only worried about the order of magnitude.
- Load
  - Requests Per Second
  - Data Volume
  - User Traffic
- Storage
  - Amount of Storage Required
- Resources:
  - Number of:
    - Servers
    - CPUs
    - Memory
- Network Bandwidth
- Latency


# Step 1: Back-of-the-Envelope Estimations for Twitter

- Assumptions:
  - 300 million monthly active users.
  - 50% of users use Twitter daily.
  - Users post 2 tweets per day, on average.
  - 10% of tweets contain media.
  - Data is stored for 5 years.
- Estimations:
  - Query per second (QPS) estimate:
    - Daily active users (DAU) =  $300 \text{ million} * 50\% = 150 \text{ million}$
    - Tweets QPS =  $150 \text{ million} * 2 \text{ tweets} / 24 \text{ hour} / 3600 \text{ seconds} = \sim 3500$
    - Peek QPS =  $2 * \text{QPS} = \sim 7000$
  - Media Storage estimate:
    - Average tweet size:
      - tweet\_id = 64 bytes
      - text = 140 bytes
      - media = 1 MB
    - Media storage:  $150 \text{ million} * 2 * 10\% * 1 \text{ MB} = 30 \text{ TB per day}$
    - 5-year media storage:  $30 \text{ TB} * 365 * 5 = \sim 55 \text{ PB}$

## Step 2: Propose High-Level Design & Get Buy-In

### 1. Sketch a diagram of your system's core functionalities

- Use the requirements you gathered in step one as a checklist



## Step 2: Propose High-Level Design & Get Buy-In

### 2. Define the API

- What endpoints will you have to create?

Endpoint	/tweet
HTTP Method	GET
Parameters	<code>integer</code> tweet_id
JSON Response	<code>integer</code> user_id <code>string</code> tweet_content

## Step 2: Propose High-Level Design & Get Buy-In


### 3. Define the Data Model

- What tables will you need to create in your database?
- What are the relationships between those tables?

tweets	
<b>tweet_id</b>	<b>int</b>
user_id	int
content	string
created_at	datetime

users	
<b>user_id</b>	<b>int</b>
username	string
email	string
created_at	datetime

likes	
<b>like_id</b>	<b>int</b>
tweet_id	int
user_id	int
created_at	datetime



## Step 2: Propose High-Level Design & Get Buy-In

### 4. Get Buy-In

- Check in with your interviewer, confirm you're on the right track
- Make space for them to interject
- Only move forward after you have buy-in from your interviewer



## Step 3: Design Deep-Dive

- Work with the interviewer to identify and prioritize components in the architecture to focus in on.
- Make the system:
  - Faster
    - CDNs, Caches
  - Robust
    - Database Replication
    - Sharding
  - Secure
    - OAuth
    - ACLs
    - Encryption


} Also makes it faster

## Step 3: Design Deep-Dive

- Technology Choices
  - e.g. What sort of database?
    - Relational vs NoSQL
  - e.g. What sort of message queue/event streaming system?
    - Kafka, RabbitMQ, etc.
  - e.g. Long polling vs Websockets?
- The choices matter, but the ability to discuss the trade offs between each choice matters more.
  - If you're going to choose NoSQL over Relational, be prepared to explain why. What might change that decision?

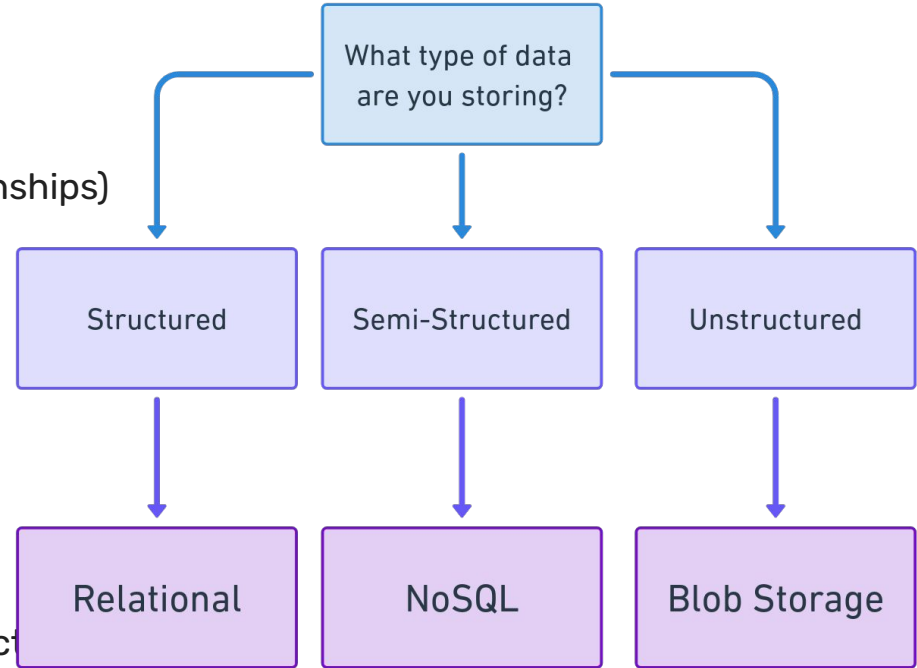
# Choice of Database System

- Choose an appropriate database technology based on the type of data you're dealing with.
- Relational is often a safe bet



# Structured, Semi-Structured, and Unstructured Data

- **Structured Data**
  - Highly organized
  - Consistent format
  - Multiple sets of data interlinked (relationships)
    - Customer Information
    - Financial Transactions
- **Semi-Structured Data**
  - Some/Simple organization
  - Variable pieces of data
    - JSON / XML
    - Log files
- **Unstructured**
  - Requires additional processing to extract
    - Images
    - Videos



## Step Four: Wrap-Up

- Give the interviewer a recap of your design.
- Discuss possible improvements, if you had more time.
- Error Cases
  - e.g. server failure
- What metrics should you track?
  - e.g. HTTP Response Codes
- How will you monitor the system?
- Make space for the interviewer to pick something to drill into.

# The Best Tip I Can Give

- Unless you're Allen Iverson, **YOU NEED PRACTICE.**
- Talk to your classmates, ask if they want to practice together
- [Pramp](#)

These slides can be found at:

[https://carloscuevas.github.io/systems\\_design.pdf](https://carloscuevas.github.io/systems_design.pdf)